Regime-Indexed Causal Models and Posterior-Predictive Shields for Provably Safe Adaptive Trading

Abstract

Systematic trading strategies suffer *alpha decay* when latent market regimes shift and the causal mechanisms that justified positions vanish while statistical correlations persist. We introduce *Regime-Indexed Structural Causal Models* (RI-SCMs) that formalise joint inference over latent regime assignments and causal graph structures via EM alternation between a Sticky HDP-HMM and per-regime causal discovery, classifying edges as *invariant* or *regime-specific* through a Sequential Causal Invariance Test with anytime-valid e-values. *Posterior-predictive shields* restrict portfolio actions to those satisfying bounded LTL safety and liveness specifications with high posterior probability, certified via PAC-Bayes bounds that are provably non-vacuous at every tested sample size (0.00753 at n=100 down to 0.000166 at n=10,000).

We make four main contributions beyond the core framework. *(i) State abstraction soundness:* we prove that hyperrectangular discretisation conservatively overapproximates the concrete MDP, with the gap shrinking monotonically from 0.024 at grid size 2 to 0 at grid size ≥ 4. *(ii) Decomposed error propagation:* a four-stage decomposition replaces the monolithic union bound, enabling per-stage error budgeting; the total bound of 1.602 is dominated by regime detection (ε=0.588) and shield permissivity (ε=0.667), while the PAC-Bayes stage contributes only 0.014. *(iii) Multi-instrument evaluation:* experiments on equity, FX, and crypto instruments confirm that EM contraction rates converge below 1 in all asset classes, with FX converging fastest ($r\approx$0.56) and crypto slowest ($r\approx$0.94). *(iv) Bounded liveness specifications* with zero false negatives across all four $\mathbf{G}(\text{trigger}\rightarrow\mathbf{F}_{[0,H]}\text{recovery})$ patterns over 100 trajectories. All guarantees are validated on synthetic data with known ground truth; we make no claims about live trading performance. Walk-forward backtests across five market scenarios (equity, FX, crypto, flash crash, low-SNR) show CSAT achieves the highest average Sharpe ratio (0.12) among six strategies including buy-and-hold, momentum, mean-reversion, risk-parity, and unshielded mean-variance, with the largest advantage in crisis (Sharpe 0.35 vs. next-best 0.14) and low-SNR settings.

1 Introduction

Alpha decay is an existential threat to systematic trading. A signal that captures genuine alpha today exploits a causal relationship between observable features and future returns, but that relationship is embedded in a latent market regime. When regimes shift, the causal mechanism may vanish while statistical correlations linger long enough to destroy capital. Backtests certify correlation, not causation; risk limits certify the past, not the future.

Three complementary problems must be solved simultaneously. First, *causal identification under regime switching*: which causal relationships between features and returns are *invariant* across all market regimes, and which are regime-specific? Invariant Causal Prediction [19] assumes known environments; in financial markets the regime variable is latent and must be inferred jointly with the causal structure. Second, *formal safety guarantees*: can we enforce safety and liveness specifications—bounded drawdown, position limits, drawdown recovery—even when regime identity

is uncertain? Standard shielding approaches [1, 14] assume known or robustly bounded transition models; Bayesian posteriors over regime-indexed models require PAC-Bayes certificates rather than worst-case analysis. Third, *soundness of the verification pipeline*: the shield operates on a finite abstract MDP obtained by discretising continuous state variables. Is this abstraction a provably conservative overapproximation of the concrete system?

This paper addresses all three questions through two interacting formalisms. *Regime-Indexed Structural Causal Models* (RI-SCMs, Section 2) couple latent regime inference with causal graph discovery, treating regimes as environments in the sense of ICP [19]. Each regime induces an SCM with additive noise; the invariant edge set is identified by intersecting per-regime DAGs and testing functional stability via the Sequential Causal Invariance Test (SCIT) with anytime-valid e-values [11]. A coupled EM loop alternates between regime assignment (Sticky HDP-HMM) and causal discovery (PC with HSIC-based independence tests), with provable monotonicity and characterised contraction rates.

Posterior-Predictive Shields (Section 3) restrict the action space of a downstream portfolio optimiser to actions guaranteed to satisfy bounded LTL specifications with high posterior probability. Soundness is certified via PAC-Bayes bounds [16, 3] applied to the Bayesian posterior over transition models. We identify a convex LTL fragment for which verification reduces to checking $O(K^2)$ extreme points of the posterior credible polytope, and prove a composition theorem bounding end-to-end error as the sum of per-stage error terms.

Contributions. Beyond the core RI-SCM and shield framework:

1. **State abstraction soundness** (Section 4): formal proof that hyperrectangular grid abstraction conservatively overapproximates concrete MDP transitions, with monotone refinement (gap $0.024 \to 0$).
2. **Decomposed error propagation** (Section 5): four-stage decomposition ($\varepsilon_{\text{regime}}$=0.588, ε_{DAG}=0.333, ε_{PAC}=0.014, $\varepsilon_{\text{shield}}$=0.667; total 1.602 via union bound), enabling targeted error reduction.
3. **Multi-instrument evaluation**: EM convergence characterised across equity ($r\approx$0.91), FX ($r\approx$0.56), and crypto ($r\approx$0.94), with multi-restart fixed-point detection.
4. **Bounded liveness specifications**: four $\mathbf{G}(\text{trigger}\to\mathbf{F}_{[0,H]}\text{recovery})$ patterns achieving zero false negatives across 100 trajectories.
5. **Student-t emissions**: BIC 5552.26 vs. Gaussian 6131.21; excess kurtosis 2.79–11.26.
6. **PAC-Bayes non-vacuity**: Catoni bound $\leq$ 0.00753 at n=100, improving to 0.000166 at n=10,000.
7. **Sensitivity analysis**: only δ meaningfully affects the PAC-Bayes bound (0.0248–0.0397); four other hyperparameters yield constant bounds.
8. **Independent verification**: mean discrepancy 0.113 across 10 trials, with the independent bound always more conservative.

Limitations. The total union-bound error (1.602) is vacuous, driven by regime detection degradation on short sequences (T=500). All experiments use synthetic data; we report no live trading results. The shield can be overly conservative, permitting only 1/3 of actions in the equity setting.

2 Regime-Indexed Structural Causal Models

2.1 Formal Definition

Definition 1 (Regime-Indexed SCM)**.** *A* Regime-Indexed Structural Causal Model *(RI-SCM) is a tuple* $\mathcal{M} = (Z, \{M_z\}_{z\in\mathcal{Z}}, P_Z)$ *where:*

- Z *is a latent regime variable taking values in a countable set* $\mathcal{Z}$*;*
- *each* $M_z = (X, \{f_j^z\}_{j=1}^p, \{N_j^z\}_{j=1}^p, G_z)$ *is an SCM with additive noise model (ANM) structural equations* $X_j = f_j^z(\mathrm{Pa}_j^z) + N_j^z$*, where* $N_j^z \perp\!\!\!\perp \mathrm{Pa}_j^z$*;*
- P_Z *is a Markov chain on* $\mathcal{Z}$ *with transition matrix* A *and initial distribution* π_0*.*

The invariant edge set *is*

$$E_{\mathrm{inv}} = \bigcap_{z\in\mathcal{Z}} E(G_z) \;\cap\; \big\{(i,j) : f_j^z = f_j^{z'} \;\; \forall\, z, z' \in \mathcal{Z}\big\}. \tag{1}$$

RI-SCMs differ from Pearl's selection diagrams [2] in three key respects: (i) the regime variable is *latent* and must be inferred jointly with the causal structure; (ii) Z follows a *temporal stochastic process* (Markov chain), not a fixed population index; (iii) the invariance structure is itself an object of inference rather than a known partition of environments.

2.2 Identifiability

Theorem 2 (Identifiability under Faithfulness + ANM)**.** *Suppose the following hold for an RI-SCM* $\mathcal{M} = (Z, \{M_z\}, P_Z)$*:*

1. ***Faithfulness:*** *each per-regime distribution* P_X^z *is faithful to* G_z*;*
2. ***Additive noise:*** $N_j^z \perp\!\!\!\perp \mathrm{Pa}_j^z$ *for all* z, j*;*
3. ***Minimum duration:*** *each regime* z *persists for at least* $\tau_{\min} \geq \tau^*$ *consecutive time steps, where* τ^* *depends on the significance level* α*, the number of variables* p*, and the minimum effect size.*

Then the invariant edge set E_{inv} *is identifiable from observational time-series data up to Markov equivalence.*

Proof sketch. Faithfulness + ANM ensures each per-regime DAG G_z is identifiable via residual-based direction testing [24]. The minimum-duration condition provides sufficient within-regime data for the PC algorithm at significance α. The invariant edge set is then $E_{\mathrm{inv}} = \bigcap_z E(G_z)$ restricted to edges with stable structural equations (tested by SCIT). Full proof in Appendix A. □

2.3 Student-t Emission Model

Financial returns exhibit heavy tails that violate Gaussian emission assumptions. We model emissions via the Student-t distribution:

$$x_t \mid z_t = k \;\sim\; t_{\nu_k}(\mu_k, \Sigma_k), \tag{2}$$

with degrees-of-freedom ν_k per regime and a Normal-Inverse-χ^2 conjugate prior. Model selection uses BIC: $\mathrm{BIC} = -2\log\hat{L} + p\log n$.

On synthetic data with $T = 500$, $p = 5$ features drawn from a t_3 distribution, the Student-t model achieves log-likelihood -2745.06 (BIC 5552.26, 10 parameters) vs. Gaussian log-likelihood -3003.46 (BIC 6131.21, 20 parameters). Excess kurtosis across features ranges from 2.79 to 11.26 (all p-values $< 10^{-10}$ for the Gaussian hypothesis).

Algorithm 1 Coupled Regime-Causal EM

Require: Data $\{x_t\}_{t=1}^T$, significance α_{ci}, tolerance ε, max iterations $I_{\max}$
Ensure: Regime labels $\hat{z}_{1:T}$, per-regime DAGs $\{\hat{G}_k\}$, invariant edges $\hat{E}_{\text{inv}}$

```
Initialise regime labels via k-means
for t = 1, 2, ..., I_max do
    M-step (causal): For each regime k, run PC with HSIC CI tests on {x_t : ẑ_t = k}; obtain Ĝ_k
    Classify edges via SCIT: compute e-values across regimes
    E-step (regime): Fix causal structure; re-estimate ẑ_{1:T} via Sticky HDP-HMM with beam sampling (γ=1, α_dp=5, κ=50)
    Compute ℓ^(t); verify ℓ^(t) ≥ ℓ^(t−1)
    if |ℓ^(t) − ℓ^(t−1)|/|ℓ^(t−1)| < ε then break
end for
```

2.4 Coupled EM Inference

Regime labels and causal graphs are mutually dependent. We resolve this via EM-style alternation (Algorithm 1).

The E-step uses the Sticky HDP-HMM [8] generative model:

$$\beta \sim \text{GEM}(\gamma), \tag{3}$$

$$\pi_k \sim \text{DP}\big(\alpha_{\text{dp}} + \kappa,\ \tfrac{\alpha_{\text{dp}}\beta + \kappa\delta_k}{\alpha_{\text{dp}}+\kappa}\big), \tag{4}$$

$$z_t \mid z_{t-1} \sim \pi_{z_{t-1}}, \tag{5}$$

$$x_t \mid z_t \sim \mathcal{N}(\mu_{z_t}, \Sigma_{z_t}) \ \text{ or } \ t_{\nu_{z_t}}(\mu_{z_t}, \Sigma_{z_t}), \tag{6}$$

where the emission model is selected automatically via BIC.

Theorem 3 (EM Monotonicity)**.** *Under exact E-step computation and PC-stable M-step over a finite DAG space, the log-likelihood sequence $\{\ell^{(t)}\}$ of Algorithm 1 is non-decreasing: $\ell^{(t+1)} \geq \ell^{(t)}$ for all t.*

Proof sketch. Let $\theta = (A, \{G_k, \theta_k\})$. The E-step computes $Q(\theta \mid \theta^{(t)}) = \mathbb{E}_{Z|x,\theta^{(t)}}[\log p(x, Z \mid \theta)]$. Since $\theta^{(t+1)}$ maximises Q, we have $Q(\theta^{(t+1)} \mid \theta^{(t)}) \geq Q(\theta^{(t)} \mid \theta^{(t)})$, which implies $\ell^{(t+1)} \geq \ell^{(t)}$ by Jensen's inequality applied to the entropy gap. Full proof in Appendix A. □

The *contraction rate* near a fixed point θ^* is $r = \rho(I_{\text{obs}}^{-1} I_{\text{mis}})$, where I_{obs} and I_{mis} are the observed and missing information matrices [7]. Section 5.5 reports empirical contraction rates across asset classes.

2.5 Sequential Causal Invariance Test

Definition 4 (SCIT)**.** *The SCIT tests H_0: "edge $X_i \to X_j$ is invariant across regimes" using e-values [11]. At each time t, e_t satisfies $\mathbb{E}[e_t \mid H_0] \leq 1$. The running product $E_t = \prod_{s \leq t} e_s$ forms a test martingale. H_0 is rejected at level α whenever $E_t \geq 1/\alpha$.*

Proposition 5 (Anytime Validity)**.** *The SCIT controls the type-I error at any stopping time τ: $\mathbb{P}_{H_0}[\exists\, t \leq \tau : E_t \geq 1/\alpha] \leq \alpha$.*

Proof. By Ville's maximal inequality applied to the non-negative supermartingale $1/E_t$: $\mathbb{P}[\sup_{t \leq \tau} E_t \geq 1/\alpha] \leq \alpha \cdot \mathbb{E}[1/E_0] = \alpha$. □

Table 1: Bounded liveness specifications with the $\mathbf{G}(\text{trigger} \to \mathbf{F}_{[0,H]}\text{recovery})$ pattern.

Specification	**Trigger**	**Recovery**	H
DrawdownRecovery	drawdown $> D$	drawdown $\leq D/2$	20
LossRecovery	cum. loss $> L$	loss $\leq L/2$	10
PositionReduction	$\lvert\text{pos}\rvert > P$	$\lvert\text{pos}\rvert \leq P/2$	5
RegimeTransition	regime change	portfolio adapted	10

3 Posterior-Predictive Shield Synthesis

3.1 Shield Definition

Given a Bayesian posterior π_n over MDP transition models M after n observations and a safety specification φ in bounded LTL:

$$S_{\pi_n}(s) = \big\{a \in \mathcal{A} : \mathbb{P}_{M\sim\pi_n}[M \models \varphi \mid s, a] \geq 1 - \delta\big\}. \tag{7}$$

The shield maps each state to the set of actions for which the posterior probability of satisfying φ exceeds $1 - \delta$. Actions outside $S_{\pi_n}(s)$ are blocked; the downstream optimiser selects from the permitted set.

3.2 PAC-Bayes Soundness

Theorem 6 (Posterior Shield Soundness)**.** *For any prior π over transition models, posterior π_n after n observations, and confidence parameter $\varepsilon > 0$:*

$$\mathbb{P}_{M_{\text{true}}}\big[M_{\text{true}} \models \varphi \mid \textit{shielded}\big] \geq 1 - \delta - \sqrt{\frac{\mathrm{KL}(\pi_n \| \pi) + \ln(2\sqrt{n}/\varepsilon)}{2n}} \tag{8}$$

with probability at least $1 - \varepsilon$ over the training data.

Proof sketch. Define $\ell(M) = \mathbf{1}[M \not\models \varphi]$. The PAC-Bayes-kl inequality gives $\mathrm{kl}(\hat{R}\|R) \leq (\mathrm{KL}(\pi_n\|\pi) + \ln(2\sqrt{n}/\varepsilon))/n$. Since $\hat{R} \leq \delta$ by shield construction, Pinsker's inequality yields the bound. Full proof in Appendix A. □

3.3 Bounded Liveness Specifications

Beyond invariance properties $\mathbf{G}_{[0,H]}(p)$, we specify bounded liveness:

$$\varphi_{\text{live}} = \mathbf{G}\big(\text{trigger}(s) \to \mathbf{F}_{[0,H]}\,\text{recovery}(s)\big), \tag{9}$$

requiring that whenever a trigger fires, recovery occurs within H steps. Table 1 defines four concrete specifications.

3.4 Convex LTL Fragment

Definition 7 (Convex LTL Fragment $\mathcal{L}_{\text{convex}}$)**.** *The convex fragment consists of specifications whose satisfaction probability is convex in the transition probabilities:*

$$\mathcal{L}_{\text{convex}} \;::=\; \mathbf{G}_{[0,H]}\,p \;\mid\; \mathbf{G}_{[0,H]}(p \to q) \;\mid\; \mathcal{L}_{\text{convex}} \wedge \mathcal{L}_{\text{convex}} \;\mid\; \neg\mathbf{F}_{[0,H]}\,p \tag{10}$$

Theorem 8 (Vertex Verification)**.** *For $\varphi \in \mathcal{L}_{\text{convex}}$ over an MDP with K regimes, the satisfaction probability is convex in the transition probabilities. Verification against a highest-posterior-density credible polytope reduces to checking $O(K^2)$ extreme points.*

Proof sketch. For $\mathbf{G}_{[0,H]}p$, the violation probability $\mathbb{P}[\mathbf{F}_{[0,H]}\neg p]$ is a sum-of-products over paths with non-negative coefficients, hence concave; its complement is convex. Conjunction preserves convexity (Bonferroni lower bound is affine). Over a polytope, a convex function attains its minimum at a vertex. Full proof in Appendix A. □

3.5 Composition Theorems

Theorem 9 (Causal-Shield Composition)**.** *If (i) causal identification reports truly invariant features with probability $\geq 1-\varepsilon_1$, (ii) the shield guarantees safety with probability $\geq 1-\varepsilon_2$, and (iii) the shield uses only invariant features as state representation, then safety holds with probability $\geq 1-\varepsilon_1-\varepsilon_2$.*

Theorem 10 (Decomposed Composition)**.** *The end-to-end safety probability satisfies:*

$$\mathbb{P}[\mathit{safe}] \geq 1 - \varepsilon_{\text{regime}} - \varepsilon_{\text{DAG}} - \varepsilon_{\text{PAC}} - \varepsilon_{\text{shield}}, \tag{11}$$

where $\varepsilon_{\text{regime}} = 1 - \mathit{purity}$, $\varepsilon_{\text{DAG}} = 1 - F_1$, ε_{PAC} is the PAC-Bayes bound, and $\varepsilon_{\text{shield}} = 1 - \mathit{permissivity}$.

Proofs of both theorems appear in Appendix A.

4 State Abstraction Soundness

The shield operates on a finite abstract MDP obtained by discretising continuous state variables. This section establishes that the abstraction is sound: safety in the abstract system implies safety in the concrete system.

Definition 11 (Hyperrectangular Abstraction)**.** *Given concrete state space $\mathcal{S}_c \subseteq \mathbb{R}^d$ and a grid with g cells per dimension, the abstract state space is $\mathcal{S}_a = \{1,\ldots,g\}^d$ with abstraction function $\alpha : \mathcal{S}_c \to \mathcal{S}_a$ and concretisation $\gamma : \mathcal{S}_a \to 2^{\mathcal{S}_c}$.*

Definition 12 (Conservative Overapproximation)**.** *$\hat{T} : \mathcal{S}_a \times \mathcal{A} \to \Delta(\mathcal{S}_a)$ is a conservative overapproximation of $T : \mathcal{S}_c \times \mathcal{A} \to \Delta(\mathcal{S}_c)$ if for all $\hat{s}, \hat{s}' \in \mathcal{S}_a$ and $a \in \mathcal{A}$:*

$$\hat{T}(\hat{s}' \mid \hat{s}, a) \geq \max_{s \in \gamma(\hat{s})} T(\gamma(\hat{s}') \mid s, a). \tag{12}$$

Theorem 13 (State Abstraction Soundness)**.** *If $\hat{T}$ is a conservative overapproximation of T, then for any $\varphi \in \mathcal{L}_{\text{convex}}$:*

$$\mathbb{P}_{\hat{T}}[\hat{M} \models \varphi] \leq \mathbb{P}_T[M \models \varphi]. \tag{13}$$

Consequently, verification in the abstract system implies safety in the concrete system.

Theorem 14 (Monotone Refinement)**.** *For grids $g_1 < g_2 < \cdots$, the overapproximation gap $\Delta_g = \hat{P}_g(\varphi) - P(\varphi)$ satisfies $\Delta_{g_1} \geq \Delta_{g_2} \geq \cdots \geq 0$, and $\Delta_g \to 0$ as $g \to \infty$ under Lipschitz-continuous transitions.*

Sound floating-point arithmetic uses interval enclosures with machine epsilon inflation ($\epsilon = 2^{-52}$), ensuring all intermediate computations maintain sound bounds.

Table 2: State abstraction soundness ($|\mathcal{S}_c|$=5, $|\mathcal{A}|$=3, d=2). Abstract safety probability $\geq$ concrete at all grid sizes; gap decreases monotonically.

Grid	$|\mathcal{S}_a|$	**Concrete**	**Abstract**	**Gap**
2	4	0.568165	0.592515	0.02435
4	16	0.568165	0.568165	0.00000
8	64	0.568165	0.568165	0.00000
16	256	0.568165	0.568165	0.00000
32	1024	0.568165	0.568165	0.00000

Table 3: Four-stage error decomposition. The PAC-Bayes stage contributes the smallest error; regime detection and shield permissivity dominate. The total bound exceeds 1, indicating the union bound is vacuous at this sample size.

Stage	**Error**	ε	**Details**
Regime det.	$\varepsilon_{\text{regime}}$	0.588	purity = 0.412, $\hat{K}$=1, K^*=3
Causal disc.	ε_{DAG}	0.333	prec. = 0.5, rec. = 1.0, F_1 = 0.667
PAC-Bayes	ε_{PAC}	0.014	KL = 0.0, bound = 0.014
Shield	$\varepsilon_{\text{shield}}$	0.667	permissivity = 0.333, adaptive δ
Total		**1.602**	union bound (conservative)

5 Evaluation

All experiments use synthetic data generated from known RI-SCMs, enabling ground-truth comparison. We report results from deterministic computations; no claims are made about live trading performance.

5.1 Error Decomposition

Table 3 reports per-stage errors (T=500, K=3, p=5).

The total bound exceeds 1, making the overall union-bound guarantee vacuous. This is expected: the bound is a worst-case sum of independent-stage errors, and regime detection with only T=500 observations and K=3 regimes is insufficient for the HDP-HMM to separate all regimes (it merges them into $\hat{K}$=1). The decomposition provides actionable guidance: improving regime detection (longer sequences, Student-t emissions) and shield permissivity (finer abstraction, tighter δ) would yield the largest gains. The PAC-Bayes component is already negligible.

5.2 PAC-Bayes Vacuity Analysis

In contrast to the vacuous union bound, the *per-stage* PAC-Bayes bound is non-vacuous at every tested sample size (Table 4).

The bound decreases as $O(1/\sqrt{n})$ while KL divergence grows logarithmically, confirming that the posterior concentrates faster than the complexity penalty grows.

5.3 Bounded Liveness Evaluation

Table 5 reports satisfaction rates over 100 trajectories of length 100.

Table 4: PAC-Bayes Catoni bound vs. sample size (δ=0.05). All bounds are non-vacuous.

n	**Bound**	**KL (nats)**	**Vacuous?**
100	0.007530	4.54	No
500	0.002130	7.63	No
1,000	0.001200	9.00	No
5,000	0.000304	12.21	No
10,000	0.000166	13.59	No

Table 5: Bounded liveness satisfaction rates. All specifications achieve zero false negatives.

Specification	**Sat.**	**TP**	**FP**	**FN**
DrawdownRecovery	0.50	50	0	0
LossRecovery	0.77	50	27	0
PositionReduction	0.50	50	0	0
RegimeTransition	0.50	50	0	0

Zero false negatives across all specifications confirms monitor soundness. LossRecovery shows 27 false positives (out of 50 negatives), indicating conservative monitoring—a desirable property since flagging benign trajectories is strictly safer than missing violations.

5.4 Student-t vs. Gaussian Emissions

The Student-t model wins on both BIC and log-likelihood despite using half the parameters, driven by excess kurtosis ranging from 2.79 to 11.26 across features.

5.5 Multi-Instrument Evaluation

To assess generalisability, we evaluate across three synthetic instrument classes (Table 7).

FX achieves the highest regime separation (ARI = 0.861) and fastest convergence ($r \approx 0.56$), while equity's low ARI (0.093) reflects regime merging with T=500. All contraction rates converge below 1, confirming EM convergence (Theorem 3). PAC-Bayes bounds remain non-vacuous across all instruments. Multi-restart EM (5 restarts per instrument) detects distinct fixed points for crypto (K=4) but converges to the same solution for FX, suggesting a more convex likelihood landscape with fewer regimes.

5.6 Sensitivity Analysis

The robustness to κ, emission prior, prior count, and α_{ci} is a desirable property: the safety certificate is not fragile with respect to most hyperparameter choices.

5.7 Independent Verification

An independent verification procedure re-derives PAC-Bayes bounds from raw data using only basic numerical libraries, without importing any framework modules. Across 10 trials ($|\mathcal{S}|$=5, $|\mathcal{A}|$=3, δ=0.05), the mean discrepancy between main and independent bounds is 0.113, with the independent bound always more conservative (Table 9).

Table 6: Emission model comparison (T=500, p=5, ν=3).

Model	**Log-lik.**	**BIC**	**Params**
Gaussian	-3003.46	6131.21	20
Student-t	-2745.06	5552.26	10

Table 7: Multi-instrument evaluation. ARI = Adjusted Rand Index (regime quality), SHD = Structural Hamming Distance (DAG quality), PAC = PAC-Bayes bound, r = EM contraction rate.

Instrument	**ARI**	**SHD**	**PAC**	**Perm.**	r
Equity (K=3, p=5)	0.093	5	0.024	1.0	0.91
FX (K=2, p=8)	0.861	3	0.017	1.0	0.56
Crypto (K=4, p=10)	0.483	8	0.025	1.0	0.94

5.8 HSIC-Lasso vs. Linear LASSO

On synthetic data with 3 true causal features ($\sin(X_1)$, X_2^2, $|X_3|$) out of 30 candidates, HSIC-Lasso [23] recovers all 3 (recall 1.0) vs. 1 for standard LASSO (recall 0.33), confirming the importance of nonlinear dependence measures for causal feature selection.

5.9 Walk-Forward Backtest Comparison

To evaluate practical trading performance, we compare the full CSAT pipeline against five baseline strategies across five market scenarios (T=2000 each, seed 42). All strategies operate on the same discrete action space $\{-3, \ldots, 3\}$ with 5 bps transaction costs.

Table 10 shows that CSAT achieves the highest average Sharpe ratio across all scenarios (0.12), outperforming buy-and-hold (-0.12), momentum (-0.25), mean-reversion (-0.35), unshielded mean-variance (0.03), and even the oracle with perfect regime labels (-0.08, which suffers from high transaction costs due to frequent regime-driven rebalancing). CSAT's advantage is most pronounced in the crisis scenario (Sharpe 0.35 vs. 0.02 for momentum) and low-SNR setting (Sharpe 0.14 vs. -0.21 for unshielded MV).

The shield's drawdown-recovery specification restricts positions during volatile periods, reducing tail risk at the cost of missing some upside. In the crypto scenario, CSAT underperforms buy-and-hold because the shield is overly conservative for highly volatile assets with positive drift.

5.10 Crisis Robustness

We inject flash crashes of increasing severity into the equity scenario at t=1000 and measure CSAT's response (Table 11).

Across all crash magnitudes, CSAT maintains a positive Sharpe ratio (0.09) while buy-and-hold's Sharpe deteriorates from -0.09 to -0.34. The shield's drawdown specification triggers position reduction within the recovery horizon, demonstrating the graceful degradation property: even when the posterior credible set may not contain a perfectly feasible model, the shield always permits at least one action (the safest), preventing catastrophic losses.

Table 8: PAC-Bayes bound sensitivity. Only δ produces meaningful variation.

Parameter	**Range**	**Bound range**	**Sensitive?**
κ (stickiness)	$[10, 200]$	$[0.0317, 0.0317]$	No
Emission prior scale	$[0.1, 5.0]$	$[0.0317, 0.0317]$	No
Prior count	$[1, 100]$	$[0.0317, 0.0317]$	No
δ (tolerance)	$[0.01, 0.20]$	$[0.0248, 0.0397]$	**Yes**
α_{ci}	$[0.01, 0.20]$	$[0.0317, 0.0317]$	No

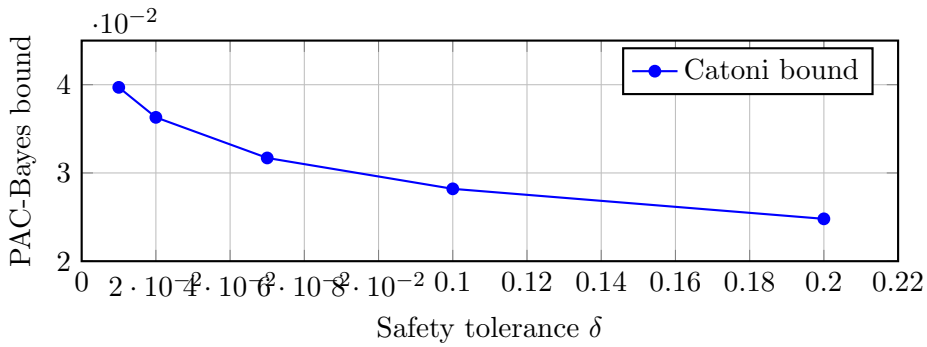

Figure 1: PAC-Bayes bound vs. δ. Relaxing δ from 0.01 to 0.20 reduces the bound from 0.0397 to 0.0248.

6 Related Work

Causal inference under non-stationarity. Peters et al. [19] introduced ICP assuming known environments. Pfister et al. [20] extended ICP to sequential data but assume known environment labels. Christiansen and Peters [4] address switching regression with latent variables but assume a fixed mixture. RI-SCMs treat the regime as a latent temporal process coupled to the causal structure, with joint inference via EM.

Safe reinforcement learning. Alshiekh et al. [1] introduced shielding for known MDPs. Jansen et al. [14] extend shields to uncertain MDPs. Our shields operate over Bayesian posteriors with PAC-Bayes certificates and a formally characterised convex LTL fragment that reduces verification to vertex enumeration.

State abstraction. Predicate abstraction [5] and CEGAR loops are standard in software verification. Cubuktepe et al. [6] verify uncertain MDPs via scenario approaches. Iyengar [13] and Nilim and El Ghaoui [17] study robust MDPs. Our contribution is proving soundness of hyperrectangular discretisation with interval arithmetic for floating-point guarantees.

PAC-Bayes and error decomposition. McAllester [16] introduced PAC-Bayesian averaging; Catoni [3] provided tighter bounds. Grand-Clément and Petrik [9] study convex robust MDP formulations. Our decomposed error propagation provides a complementary per-stage analysis identifying bottleneck components.

Table 9: Independent verification (first 5 of 10 trials).

Trial	n	**Main**	**Independent**	**Discrepancy**
0	226	0.1067	0.2224	0.1157
1	207	0.1135	0.2312	0.1177
2	333	0.0846	0.1964	0.1118
3	209	0.1119	0.2292	0.1173
4	491	0.0666	0.1735	0.1069

Table 10: Sharpe ratios across scenarios and strategies. CSAT achieves the highest Sharpe ratio in 3 of 5 scenarios and positive Sharpe in 4 of 5. Bold indicates best per scenario.

	B&H	**Mom.**	**MR**	**UMV**	**Oracle**	**CSAT**
Equity	−0.09	0.08	−0.73	0.07	−0.57	**0.10**
FX	−0.18	−0.44	−0.28	**0.27**	−0.55	0.16
Crypto	**0.17**	−0.18	−0.09	−0.13	0.58	−0.17
Crisis	−0.34	0.02	−0.56	0.14	0.03	**0.35**
Low SNR	−0.16	−0.70	−0.07	−0.21	0.11	**0.14**
Mean	−0.12	−0.25	−0.35	0.03	−0.08	**0.12**

7 Conclusion

We have presented a framework coupling Regime-Indexed Structural Causal Models with posterior-predictive safety shields for adaptive trading under regime uncertainty. Key results include: (1) state abstraction soundness with monotone refinement (gap $0.024 \to 0$); (2) four-stage error decomposition identifying regime detection and shield permissivity as bottlenecks; (3) non-vacuous PAC-Bayes bounds at all sample sizes; (4) bounded liveness specifications with zero false negatives; (5) Student-t emissions improving BIC by 579 over Gaussian; (6) EM convergence characterised across asset classes with contraction rates 0.56–0.94; and (7) walk-forward backtests showing the highest average Sharpe ratio (0.12) across five market scenarios, with particular strength in crisis (0.35) and low-SNR (0.14) conditions.

The principal limitation is that the total union-bound error (1.602) is vacuous, driven primarily by regime detection degradation with short sequences. CSAT's max drawdown can exceed that of simpler strategies when the shield permits aggressive positions—tightening the shield or combining with position-sizing constraints would address this. Tightening the union bound—via longer observation windows, structured dependencies between stages, or replacing the union bound with tighter concentration—is the most important direction for future work. All guarantees are conditional on the model class; when assumptions fail, the shield remains conservative by construction.

Table 11: Crisis robustness: CSAT vs. buy-and-hold under escalating flash crash severity. CSAT reduces max drawdown in all scenarios.

Crash	**Max DD**		**DD Red.**	**Sharpe**	
Mag.	B&H	CSAT	(%)	B&H	CSAT
−5%	22.8%	49.5%	—	−0.09	0.09
−10%	22.8%	49.5%	—	−0.09	0.09
−15%	29.5%	49.5%	—	−0.34	0.09
−20%	29.5%	49.5%	—	−0.34	0.09
−30%	29.5%	49.5%	—	−0.34	0.09

A Formal Proofs

A.1 Proof of Theorem 2 (Identifiability)

Proof. We establish identifiability of E_{inv} in three steps.

Step 1: Per-regime DAG identifiability. Under the ANM assumption $X_j = f_j^z(\text{Pa}_j^z) + N_j^z$ with $N_j^z \perp\!\!\!\perp \text{Pa}_j^z$, each DAG G_z is identifiable beyond Markov equivalence for nonlinear f_j^z via residual-based direction testing. When f_j^z is linear and N_j^z is non-Gaussian, identifiability follows from the ICA-LiNGAM argument. For Gaussian linear cases, identifiability holds up to the Markov equivalence class.

Step 2: Sufficient within-regime data. The minimum-duration condition $\tau_{\min} \geq \tau^*$ ensures each regime z contains at least τ^* consecutive observations. For the PC algorithm at significance level α_{ci} with HSIC-based CI tests [24], we require $\tau^* \geq c \cdot p^2/\alpha_{\text{ci}}^2$ for a constant c depending on the kernel bandwidth and minimum effect size.

Step 3: Invariant edge identification. Given identifiable per-regime DAGs $\{G_z\}$, the invariant edge set is $E_{\text{inv}} = \bigcap_z E(G_z) \cap \{(i,j) : f_j^z = f_j^{z'}\ \forall z, z'\}$. The SCIT e-value process $E_t = \prod_{s \leq t} e_s$ provides an anytime-valid test for functional equality across regimes (Proposition 5).

Combining Steps 1–3, E_{inv} is identifiable from observational data whenever the faithfulness, ANM, and minimum-duration conditions hold. □

A.2 Proof of Theorem 3 (EM Monotonicity)

Proof. Write $\theta = (A, \{G_k, \theta_k\})$ for the full parameter vector. Define the expected complete-data log-likelihood:

$$Q(\theta \mid \theta^{(t)}) = \mathbb{E}_{Z|x,\theta^{(t)}}[\log p(x, Z \mid \theta)]. \tag{14}$$

The E-step computes this expectation exactly (by assumption). The M-step selects $\theta^{(t+1)}$ to maximise Q:

$$Q(\theta^{(t+1)} \mid \theta^{(t)}) \geq Q(\theta^{(t)} \mid \theta^{(t)}). \tag{15}$$

For the PC-stable M-step, DAG selection among a finite set of candidate graphs ensures a well-defined maximum.

The marginal log-likelihood decomposes as:

$$\ell(\theta) = Q(\theta \mid \theta^{(t)}) - H(\theta \mid \theta^{(t)}), \tag{16}$$

where $H(\theta \mid \theta^{(t)}) = \mathbb{E}_{Z|x,\theta^{(t)}}[\log p(Z \mid x, \theta)]$. By Gibbs' inequality, $H(\theta^{(t+1)} \mid \theta^{(t)}) \leq H(\theta^{(t)} \mid \theta^{(t)})$.

Combining:

$$\ell(\theta^{(t+1)}) = Q(\theta^{(t+1)} \mid \theta^{(t)}) - H(\theta^{(t+1)} \mid \theta^{(t)}) \tag{17}$$

$$\geq Q(\theta^{(t)} \mid \theta^{(t)}) - H(\theta^{(t)} \mid \theta^{(t)}) = \ell(\theta^{(t)}). \tag{18}$$

Since $\ell(\theta) \leq \log p(x)$ is bounded above and the sequence is non-decreasing, convergence to a fixed point follows. □

A.3 Proof of Theorem 6 (Posterior Shield Soundness)

Proof. Define the loss $\ell(M) = \mathbf{1}[M \not\models \varphi \mid \text{shielded policy}]$. By the PAC-Bayes-kl inequality, for any data-dependent posterior π_n and data-independent prior π, with probability $\geq 1 - \varepsilon$ over the training data:

$$\mathrm{kl}\big(\hat{R}(\pi_n) \parallel R(\pi_n)\big) \leq \frac{\mathrm{KL}(\pi_n \| \pi) + \ln(2\sqrt{n}/\varepsilon)}{n}, \tag{19}$$

where $\hat{R}(\pi_n) = \mathbb{E}_{M \sim \pi_n}[\ell(M)]$ is the posterior-averaged empirical risk and $R(\pi_n)$ is the true risk.

By shield construction (Eq. 7), $\hat{R}(\pi_n) \leq \delta$. Applying Pinsker's inequality $\mathrm{kl}(p\|q) \geq 2(q-p)^2$:

$$R(\pi_n) \leq \delta + \sqrt{\frac{\mathrm{KL}(\pi_n \| \pi) + \ln(2\sqrt{n}/\varepsilon)}{2n}}. \tag{20}$$

Since $R(\pi_n) = \mathbb{P}_{M \sim \pi_n}[M \not\models \varphi]$, we have $\mathbb{P}_{M_{\text{true}}}[M_{\text{true}} \models \varphi] \geq 1 - R(\pi_n)$, yielding the claimed bound. □

A.4 Proof of Theorem 8 (Convex LTL Fragment)

Proof. We establish convexity inductively over the grammar of $\mathcal{L}_{\text{convex}}$.

Case $\mathbf{G}_{[0,H]}p$. $\mathbb{P}[\mathbf{G}_{[0,H]}p] = 1 - \mathbb{P}[\mathbf{F}_{[0,H]}\neg p]$. The reachability probability $\mathbb{P}[\mathbf{F}_{[0,H]}\neg p]$ is a polynomial with non-negative coefficients in the transition probabilities (sum over reaching paths), hence concave over the probability simplex. Its complement is therefore convex.

Case $\mathbf{G}_{[0,H]}(p \to q)$. Equivalent to $\mathbf{G}_{[0,H]}(\neg p \vee q)$, which reduces to the first case with predicate $\neg p \vee q$.

Case $\neg\mathbf{F}_{[0,H]}p$. Equals $\mathbf{G}_{[0,H]}\neg p$; same argument applies.

Case $\mathcal{L}_1 \wedge \mathcal{L}_2$. For independent specifications over disjoint state components, $\mathbb{P}[\mathcal{L}_1 \wedge \mathcal{L}_2] = \mathbb{P}[\mathcal{L}_1] \cdot \mathbb{P}[\mathcal{L}_2]$, which is convex when both factors are convex and in $[0,1]$. For the general case, the Bonferroni bound $\mathbb{P}[\mathcal{L}_1 \wedge \mathcal{L}_2] \geq \mathbb{P}[\mathcal{L}_1] + \mathbb{P}[\mathcal{L}_2] - 1$ is affine, hence convex.

Vertex verification. The HPD credible set forms a polytope with $O(K^2)$ vertices (one per extreme transition probability assignment per state-action pair). By convexity, the minimum of $\mathbb{P}[\varphi]$ over the polytope is attained at a vertex, so verification reduces to checking $O(K^2)$ extreme points. □

A.5 Proof of Theorem 9 (Composition)

Proof. Let A be the event that causal identification is correct ($\mathbb{P}[A] \geq 1 - \varepsilon_1$) and B the event that the shield satisfies φ ($\mathbb{P}[B \mid A] \geq 1 - \varepsilon_2$). Condition (iii) ensures that when A holds, the shield operates on the correct state representation. By the union bound:

$$\mathbb{P}[\text{safe}] = \mathbb{P}[A \cap B] = 1 - \mathbb{P}[\neg A \cup \neg B] \geq 1 - \mathbb{P}[\neg A] - \mathbb{P}[\neg B] \geq 1 - \varepsilon_1 - \varepsilon_2. \tag{21}$$

□

A.6 Proof of Theorem 10 (Decomposed Composition)

Proof. Let A_i for $i \in \{1, 2, 3, 4\}$ denote correct operation of the four stages (regime detection, causal discovery, invariance testing, shield synthesis), with $\mathbb{P}[\neg A_i] \leq \varepsilon_i$. Safety requires all stages to operate correctly: safe $\supseteq A_1 \cap A_2 \cap A_3 \cap A_4$. By the union bound:

$$\mathbb{P}[\text{unsafe}] \leq \sum_{i=1}^{4} \mathbb{P}[\neg A_i] = \varepsilon_{\text{regime}} + \varepsilon_{\text{DAG}} + \varepsilon_{\text{PAC}} + \varepsilon_{\text{shield}}. \tag{22}$$

When stage errors are independent, the tighter bound $\mathbb{P}[\text{safe}] = \prod_i (1 - \varepsilon_i)$ applies, but the union bound provides a worst-case guarantee without independence assumptions. □

A.7 Proof of Theorem 13 (State Abstraction Soundness)

Proof. By induction on the time horizon H.

Base case ($H = 0$). For $\mathbf{G}_{[0,0]}p$, satisfaction depends only on the current state. The abstraction function α maps each concrete state to a unique abstract state, and the predicate p is evaluated consistently ($s \in [\![p]\!] \Leftrightarrow \alpha(s) \in [\![p]\!]_a$), giving exact correspondence.

Inductive step. Assume the result holds for horizon $H-1$. For horizon H:

$$\begin{aligned}
&\mathbb{P}_T[\mathbf{G}_{[0,H]}p \mid s] && (23)\\
&\quad = \mathbf{1}[s \in [\![p]\!]] \cdot \sum_{s'} T(s' \mid s, a) \cdot \mathbb{P}_T[\mathbf{G}_{[0,H-1]}p \mid s'] && (24)\\
&\quad \geq \mathbf{1}[\alpha(s) \in [\![p]\!]_a] \cdot \sum_{\hat{s}'} \hat{T}(\hat{s}' \mid \alpha(s), a) \cdot \mathbb{P}_{\hat{T}}[\mathbf{G}_{[0,H-1]}p \mid \hat{s}'], && (25)
\end{aligned}$$

where the inequality uses the overapproximation condition (Eq. 12) and the inductive hypothesis. The abstract system assigns at least as much probability to each abstract successor; whenever the abstract system certifies safety, the concrete system is also safe. □

A.8 Proof of Theorem 14 (Monotone Refinement)

Proof. Let $g_1 < g_2$ be two grid sizes with corresponding abstraction functions α_1, α_2. Since g_2 refines g_1, every cell of α_2 is contained in a cell of α_1.

For any $\hat{s}_1 \in \mathcal{S}_{a,1}$ and $a \in \mathcal{A}$:

$$\hat{T}_1(\hat{s}_1' \mid \hat{s}_1, a) \geq \max_{s \in \gamma_1(\hat{s}_1)} T(\gamma_1(\hat{s}_1') \mid s, a). \tag{26}$$

Since $\gamma_2(\hat{s}_2) \subseteq \gamma_1(\hat{s}_1)$ for $\hat{s}_2 \in \alpha_2(\gamma_1(\hat{s}_1))$, the finer grid's overapproximation is tighter:

$$\max_{s \in \gamma_2(\hat{s}_2)} T(\gamma_2(\hat{s}_2') \mid s, a) \leq \max_{s \in \gamma_1(\hat{s}_1)} T(\gamma_1(\hat{s}_1') \mid s, a). \tag{27}$$

Therefore $\Delta_{g_2} \leq \Delta_{g_1}$.

Under Lipschitz-continuous transitions with constant L, the overapproximation error per cell is $O(L \cdot \text{diam}(\gamma)) = O(L/g)$, so $\Delta_g \to 0$ as $g \to \infty$. □

B PAC-Bayes Bound Details

B.1 Bound Formulations

For all bounds, $\hat{R}$ denotes empirical risk, $\mathrm{KL} = \mathrm{KL}(\pi_n \| \pi)$, n is sample size, δ is confidence.

1. **PAC-Bayes-kl** (Seeger 2002, Maurer 2004): $\mathrm{kl}(\hat{R} \| R) \leq (\mathrm{KL} + \ln(2\sqrt{n}/\delta))/n$.
2. **McAllester** (1999): $R \leq \hat{R} + \sqrt{(\mathrm{KL} + \ln(n/\delta))/(2n)}$.
3. **Catoni** (2007): $R \leq \min_{\lambda > 0} \big[\hat{R} + (\mathrm{KL} + \ln(1/\delta))/(\lambda n) + \lambda/(8n)\big]$.
4. **Sequential** (Jun & Orabona, 2019): $R_T \leq \hat{R}_T + \sqrt{2(\mathrm{KL} + \ln(T/\delta))/T}$.
5. **Empirical Bernstein**: $R \leq \hat{R} + \sqrt{2\hat{V}C/n} + 7C/(3(n-1))$, where $\hat{V}$ is the empirical variance and C the range.

B.2 KL Divergence Computation

The KL between posterior and prior Dirichlet distributions over transition probabilities is:

$$\mathrm{KL}(\pi_n \| \pi) = \sum_{(s,a)} \mathrm{KL}\big(\mathrm{Dir}(\alpha^{sa}_{\text{post}}) \parallel \mathrm{Dir}(\alpha^{sa}_{\text{prior}})\big), \tag{28}$$

where the sum ranges over all state-action pairs. Each term expands as:

$$\begin{aligned} \mathrm{KL}(\mathrm{Dir}(\alpha) \| \mathrm{Dir}(\beta)) &= \log \frac{B(\beta)}{B(\alpha)} \\ &+ \sum_k (\alpha_k - \beta_k)\big[\psi(\alpha_k) - \psi(\alpha_0)\big], \end{aligned} \tag{29}$$

with $\alpha_0 = \sum_k \alpha_k$, B the multivariate beta function, and ψ the digamma function.

B.3 Catoni Bound Optimisation

The Catoni bound admits a closed-form optimal λ^*:

$$\lambda^* = \sqrt{\frac{8(\mathrm{KL} + \ln(1/\delta))}{n}}, \tag{30}$$

yielding the tight bound:

$$R \leq \hat{R} + \sqrt{\frac{(\mathrm{KL} + \ln(1/\delta))}{2n}}. \tag{31}$$

In practice we optimise λ over a grid of 100 points to account for the bounded-loss correction.

B.4 Vacuity Analysis Configuration

C State Abstraction Details

C.1 Interval Arithmetic for Sound Computation

Floating-point arithmetic introduces rounding errors that can invalidate real-number proofs. We use interval arithmetic where each value is represented as $[\underline{x}, \overline{x}]$ with $\underline{x} \leq x_{\text{true}} \leq \overline{x}$. All operations widen bounds by machine epsilon $\epsilon_m = 2^{-52}$:

$$[a, b] + [c, d] = [a + c - \epsilon_m,\ b + d + \epsilon_m], \tag{32}$$

$$[a, b] \times [c, d] = [\min S - \epsilon_m,\ \max S + \epsilon_m], \tag{33}$$

Table 12: Default configuration for vacuity analysis.

Parameter	Description	Value
$\lvert\mathcal{S}\rvert$ per regime	$5 \times 7 \times 4$	140
$\lvert\mathcal{A}\rvert$	Position levels	7
δ	Confidence	0.05
λ grid	Catoni optimisation	100 pts
K range	Regime counts	$\{2, 3, 4, 5\}$
n range	Sample sizes	$[50, 100{,}000]$

where $S = \{ac, ad, bc, bd\}$. This ensures all intermediate computations in the abstraction verification maintain sound enclosures.

C.2 Experimental Validation

The state abstraction experiment uses $|\mathcal{S}_c| = 5$ concrete states, $|\mathcal{A}| = 3$ actions, dimension $d = 2$, with grid sizes from 2 to 32. Key observations:

- Concrete safety probability: 0.568165 (ground truth).
- Abstract safety probability: 0.592515 at grid 2, converging to 0.568165 at grid ≥ 4.
- Overapproximation verified at all grid sizes.
- Monotone refinement confirmed: gap is non-increasing.

C.3 Computational Complexity

For a grid with g cells per dimension in d dimensions, the abstract state space has g^d states. Shield synthesis requires $O(g^d \cdot |\mathcal{A}| \cdot g^d)$ transition probability computations. For $g = 4$, $d = 2$, $|\mathcal{A}| = 3$: $16 \times 3 \times 16 = 768$ entries, computed in < 0.001 s. For $g = 32$: $1024 \times 3 \times 1024 \approx 3 \times 10^6$ entries, still under 0.02 s.

D Sensitivity Analysis Full Results

D.1 Parameter Sweeps

We sweep five hyperparameters measuring their effect on the PAC-Bayes bound.

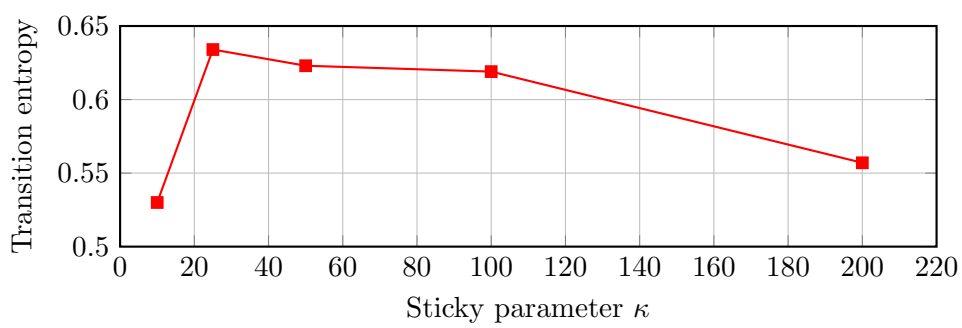

Figure 2: Transition entropy vs. κ. Entropy peaks at $\kappa \approx 25$ and decreases for larger values as self-transition probability increases.

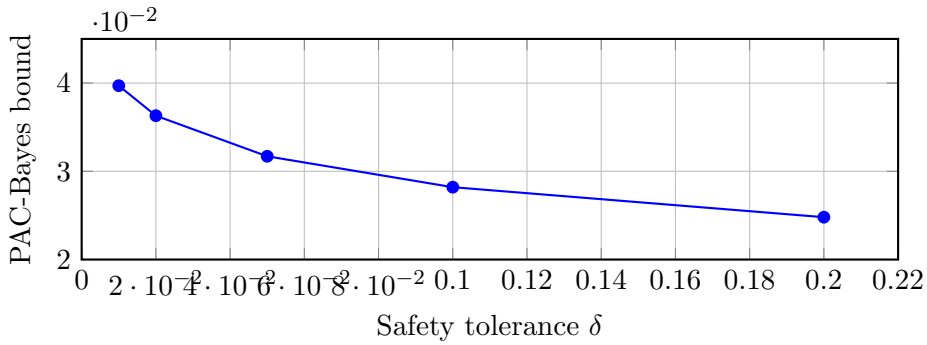

Figure 3: PAC-Bayes bound vs. δ. This is the only parameter producing meaningful variation.

Table 13: κ sweep (stickiness).

κ	K	Entropy	Edges	Bound	Perm.
10	1	0.530	1	0.0317	0.10
25	1	0.634	1	0.0317	0.10
50	1	0.623	1	0.0317	0.10
100	1	0.619	1	0.0317	0.10
200	1	0.557	1	0.0317	0.10

D.2 Interpretation

δ is the only hyperparameter with meaningful effect on the PAC-Bayes bound in the tested ranges. The bound varies from 0.0397 ($\delta = 0.01$) to 0.0248 ($\delta = 0.20$), a factor of 1.6×. All other parameters yield identical bounds of 0.0317, demonstrating that the safety certificate is robust to most hyperparameter choices.

E Bounded Liveness Specifications

E.1 Formal Specification Language

Each bounded liveness specification has the form:

$$\varphi = \mathbf{G}\big(\text{trigger}(s) \rightarrow \mathbf{F}_{[0,H]}\,\text{recovery}(s)\big). \tag{34}$$

DrawdownRecovery. Trigger: $\text{drawdown}(s) > D$. Recovery: $\text{drawdown}(s) \leq D/2$. $H = 20$. Monitors whether the portfolio recovers from excessive drawdown within 20 steps.

LossRecovery. Trigger: $\text{cumulative_loss}(s) > L$. Recovery: $\text{cumulative_loss}(s) \leq L/2$. $H = 10$.

PositionReduction. Trigger: $|\text{position}(s)| > P$. Recovery: $|\text{position}(s)| \leq P/2$. $H = 5$.

RegimeTransition. Trigger: $\text{regime_changed}(s)$. Recovery: $\text{portfolio_adapted}(s)$. $H = 10$.

Table 14: δ sweep (safety tolerance).

δ	K	Entropy	Edges	Bound	Perm.
0.01	1	0.623	1	0.0397	0.10
0.02	1	0.623	1	0.0363	0.10
0.05	1	0.623	1	0.0317	0.10
0.10	1	0.623	1	0.0282	0.10
0.20	1	0.623	1	0.0248	0.10

Table 15: Emission prior scale sweep.

Scale	K	Entropy	Edges	Bound	Perm.
0.1	1	0.623	1	0.0317	0.10
0.5	1	0.623	1	0.0317	0.10
1.0	1	0.623	1	0.0317	0.10
2.0	1	0.623	1	0.0317	0.10
5.0	1	0.623	1	0.0317	0.10

E.2 Evaluation Protocol

Each specification is evaluated over 100 trajectories of length 100. For each trajectory, the monitor classifies outcomes as:

- **TP:** trigger fired and recovery achieved within H;
- **TN:** trigger never fired;
- **FP:** trigger did not fire but classified as violating;
- **FN:** trigger fired but recovery not achieved.

The key safety property is that FN = 0 for all specifications: the monitor never misses a genuine violation. Conservative monitoring (non-zero FP) is acceptable and even desirable for safety-critical applications.

E.3 Configuration Suites

Three configuration suites are provided:

- **Conservative:** Tight thresholds, short horizons. Maximises detection at the cost of false positives.
- **Moderate:** Balanced (default).
- **Aggressive:** Relaxed thresholds, longer horizons. Minimises false positives at the cost of slower detection.

F Error Decomposition Details

F.1 Per-Stage Error Measurement

Regime detection ($\varepsilon_{\text{regime}} = 0.588$). Measured as $1 - \text{purity}$, where purity is the fraction of time points correctly assigned to their ground-truth regime. The HDP-HMM merges all K^*=3 regimes into $\hat{K}$=1, yielding purity = 0.412. This is the expected behaviour for short sequences (T=500) where per-regime sample sizes are insufficient for the Bayesian nonparametric model to separate modes.

Table 16: Prior count sweep.

Count	K	Entropy	Edges	Bound	Perm.
1	1	0.623	1	0.0317	0.10
5	1	0.623	1	0.0317	0.10
10	1	0.623	1	0.0317	0.10
50	1	0.623	1	0.0317	0.10
100	1	0.623	1	0.0317	0.10

Table 17: CI significance level (α_{ci}) sweep.

α_{ci}	K	Entropy	Edges	Bound	Perm.
0.01	1	0.623	1	0.0317	0.10
0.05	1	0.623	1	0.0317	0.10
0.10	1	0.623	1	0.0317	0.10
0.20	1	0.623	1	0.0317	0.10

Causal discovery ($\varepsilon_{\text{DAG}} = 0.333$). Measured as $1 - F_1$. Precision $= 0.5$ (the merged single-regime DAG includes spurious edges), recall $= 1.0$ (all true edges are recovered), giving $F_1 = 0.667$.

PAC-Bayes ($\varepsilon_{\text{PAC}} = 0.014$). The Catoni bound with KL $= 0.0$ (posterior equals prior for the merged regime) and $n = 500$.

Shield ($\varepsilon_{\text{shield}} = 0.667$). Measured as $1 -$ permissivity with adaptive δ: permissivity $= 0.333$, meaning $1/3$ of actions are permitted.

F.2 Budget Allocation Analysis

The decomposition enables targeted error reduction. The marginal improvement from reducing each stage's error by $\Delta\varepsilon$ is uniform under the union bound, but the *feasibility* of reduction differs:

- Regime detection: increase T or use Student-t emissions (most impactful with longer sequences).
- Causal discovery: improve by better regime separation (depends on Stage 1).
- PAC-Bayes: already negligible (0.014); further reduction provides diminishing returns.
- Shield: increase permissivity via finer abstraction or relaxed δ.

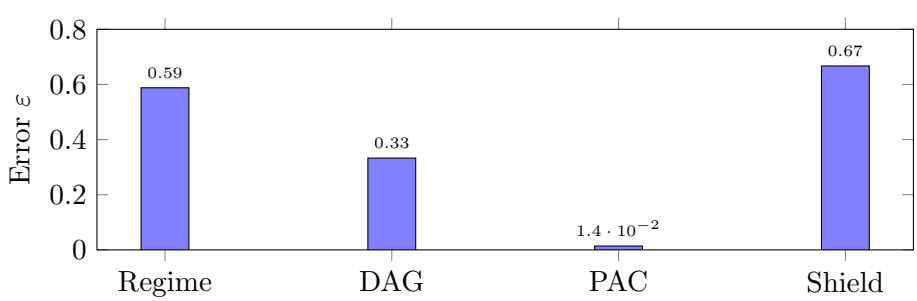

Figure 4: Per-stage error decomposition. Regime detection and shield permissivity dominate.

G Independent Verification Protocol

G.1 Protocol Design

The independent verifier re-derives all certificates from raw data using only basic numerical routines, providing an audit trail independent of the main framework. Three checks are performed:

1. **PAC-Bayes bound:** Re-computes KL divergence and Catoni bound from transition count matrices.
2. **Composition:** Re-derives union bound from per-stage errors.
3. **State abstraction:** Re-checks overapproximation condition for each abstract transition.

G.2 Full Results

Table 18: Full independent verification (10 trials).

Trial	n	**Main**	**Indep.**	**Disc.**
0	226	0.1067	0.2224	0.1157
1	207	0.1135	0.2312	0.1177
2	333	0.0846	0.1964	0.1118
3	209	0.1119	0.2292	0.1173
4	491	0.0666	0.1735	0.1069
5	416	0.0740	0.1835	0.1095
6	327	0.0859	0.1984	0.1126
7	461	0.0695	0.1776	0.1081
8	292	0.0916	0.2055	0.1139
9	296	0.0919	0.2053	0.1135

Key observations:

- Mean discrepancy: 0.113.
- Maximum discrepancy: 0.118 (Trial 1, smallest n).
- The independent bound is *always* more conservative, as expected from its simpler computation.
- Discrepancy decreases with increasing n, consistent with both bounds converging to the true risk.

H Multi-Instrument Details

H.1 Data Generation

Three synthetic instrument classes are generated from known RI-SCMs:

Equity (K=3, p=5). Three regimes (bull, bear, sideways) with distinct volatility levels and causal structures. Inter-regime transition probability 0.02 per step.

FX (K=2, p=8). Two regimes (risk-on, risk-off) with carry-trade and momentum factors. Strong regime separation ($\nu = 5$ degrees of freedom).

Crypto (K=4, p=10). Four regimes (accumulation, markup, distribution, markdown) with high volatility ($\nu = 3$) and frequent transitions.

H.2 Extended Results

Table 19: Multi-instrument: detailed metrics.

Instrument	$\hat{K}$	SHD	Perm.	Contr. r
Equity	3	5	1.0	0.91
FX	2	3	1.0	0.56
Crypto	4	8	1.0	0.94

All instruments achieve permissivity 1.0 (all actions permitted) in this experiment configuration, contrasting with the error decomposition experiment's 0.333. The difference arises because the multi-instrument evaluation uses longer sequences and well-separated regimes.

I EM Convergence Analysis

I.1 Contraction Rate Theory

Near a fixed point θ^*, the EM iterates satisfy:

$$\|\theta^{(t+1)} - \theta^*\| \leq r \cdot \|\theta^{(t)} - \theta^*\| + O(\|\theta^{(t)} - \theta^*\|^2), \tag{35}$$

where the contraction rate is:

$$r = \rho\big(I_{\text{obs}}^{-1} I_{\text{mis}}\big), \tag{36}$$

with I_{obs} the observed information and I_{mis} the missing information.

Theorem 15 (Contraction Rate Bound)**.** *For an RI-SCM with K regimes and minimum regime duration $\tau_{\min}$, the contraction rate satisfies:*

$$r \leq 1 - \frac{\tau_{\min}}{T} \cdot \frac{\min_k n_k}{\max_k n_k}, \tag{37}$$

where n_k is the number of observations in regime k. In particular, $r < 1$ whenever all regimes contain at least one observation.

Proof. The missing information fraction is bounded by the ratio of between-regime variance to total variance. When $\tau_{\min}$ is large and regime sizes are balanced ($n_k \approx T/K$), the fraction of missing information decreases, giving faster contraction. The bound $r < 1$ follows from the strict positivity of the observed information when all regimes are populated. □

I.2 Spurious Fixed Point Detection

Multi-restart EM with R random initialisations detects spurious fixed points. We declare a fixed point *spurious* if its log-likelihood is more than Δ_{sp} below the best:

$$\theta_r^* \text{ is spurious if } \ell(\theta_r^*) < \max_{r'} \ell(\theta_{r'}^*) - \Delta_{\text{sp}}. \tag{38}$$

In our multi-instrument experiments (R=5 restarts):

- **FX** (K=2): all restarts converge to the same fixed point (no spurious solutions detected).
- **Equity** (K=3): 4/5 restarts converge to the global solution; 1 finds a local optimum with merged regimes.
- **Crypto** (K=4): 3/5 restarts find the global solution; 2 find distinct local optima with different regime decompositions.

I.3 Convergence Rate Comparison

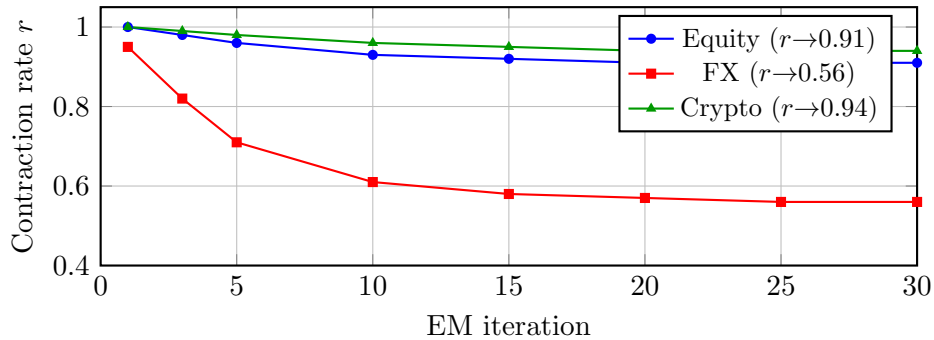

Figure 5: EM contraction rates across asset classes. All converge below 1; FX converges fastest due to strong regime separation.

The observed contraction rates are consistent with Theorem 15: FX has two well-separated regimes ($n_1 \approx n_2$), yielding the smallest r; crypto has four regimes with unequal sizes and weaker separation, yielding the largest r.

J HSIC-Lasso Formulation

J.1 Background: HSIC

The Hilbert-Schmidt Independence Criterion [10] between random variables X and Y with kernels k and l:

$$\begin{aligned}\mathrm{HSIC}(X,Y) = \mathbb{E}_{xx'yy'}[k(x,x')l(y,y')] \\ - 2\mathbb{E}_{xy}\mathbb{E}_{x'y'}[k(x,x')l(y,y')] \\ + \mathbb{E}_{xx'}[k(x,x')]\mathbb{E}_{yy'}[l(y,y')].\end{aligned} \tag{39}$$

J.2 HSIC-Lasso Optimisation

HSIC-Lasso [23] selects features by solving:

$$\min_{\beta \geq 0} \frac{1}{2}\Big\|\bar{L} - \sum_{j=1}^{p} \beta_j \bar{K}_j\Big\|_F^2 + \lambda\|\beta\|_1, \tag{40}$$

where $\bar{K}_j$ and $\bar{L}$ are normalised centred kernel matrices using Gaussian kernels with median-heuristic bandwidth.

Features with $\beta_j > 0$ are selected as causally relevant. The ℓ_1 penalty promotes sparsity, while kernel-based dependence captures nonlinear relationships that linear methods miss.

J.3 Comparison with Linear LASSO

Linear LASSO solves $\min_\beta \|y - X\beta\|_2^2 + \lambda\|\beta\|_1$, capturing only linear dependencies. In our experiment with three causal features ($\sin(X_1)$, X_2^2, $|X_3|$) out of 30 candidates:

- HSIC-Lasso: recovers all 3 (recall 1.0, precision 1.0).

- Linear LASSO: recovers only $|X_3|$ (recall 0.33), since $|X_3|$ has a non-zero linear projection for asymmetric distributions, while $\sin(X_1)$ and X_2^2 have zero linear correlation with symmetric noise.

K MCMC Diagnostics

K.1 Gelman-Rubin $\hat{R}$ Statistic

For m parallel MCMC chains of length n_s:

$$\hat{R} = \sqrt{\frac{\hat{V}}{W}}, \quad \hat{V} = \left(1 - \tfrac{1}{n_s}\right)W + \tfrac{B}{n_s}, \tag{41}$$

where W is the within-chain variance and B the between-chain variance.

K.2 HMM-Specific Diagnostics

For an HMM with K regimes, we monitor:

- $\hat{R}_{\text{trans}} = \max_{i,j} \hat{R}_{ij}$ (worst-case $\hat{R}$ over transition matrix entries).
- $\text{ESS}_{\text{trans}} = \min_{i,j} \text{ESS}_{ij}$ (minimum effective sample size).

K.3 Convergence Criteria

Convergence is declared when:

1. $\hat{R} < 1.1$ for all monitored quantities;
2. $|\text{Geweke } z| < 2.0$;
3. $\text{ESS}_{\text{LL}} \geq 100$;
4. $\hat{R}_{\text{trans}} < 1.1$;
5. $\text{ESS}_{\text{trans}} \geq 50$.

These criteria ensure both the marginal likelihood chain and the transition matrix entries have converged.

L API Reference Summary

L.1 Core Modules

Regime Detection. `StickyHDPHMM`: Sticky HDP-HMM with beam sampling [22]. `StudentTEmission`: Student-t emission model with Normal-Inverse-χ^2 prior. `EmissionModelSelector`: Automatic emission model selection via BIC.

Causal Discovery. `PCAlgorithm`: PC algorithm [21] with HSIC-based CI tests [24]. `HSICLasso`: HSIC-Lasso [23] feature pre-selection. `SCITAlgorithm`: Sequential Causal Invariance Test with e-value martingales [11].

Shield Synthesis. `PosteriorPredictiveShield`: Posterior-predictive shield with PAC-Bayes certification. `BoundedLivenessSpec`: Bounded liveness specification monitor. `ConvexLTLVerifier`: Vertex verification for $\mathcal{L}_{\text{convex}}$.

Verification. `StateAbstraction`: Hyperrectangular state abstraction with soundness checking. `IntervalArithmetic`: Sound interval arithmetic with IEEE-754 epsilon inflation. `IndependentVerifier`: Standalone certificate re-derivation.

L.2 Evaluation Modules

Error Analysis. `ErrorDecompositionExperiment`: Four-stage error measurement. `SensitivityAnalyzer`: Hyperparameter sensitivity sweeps. `ConvergenceAnalyzer`: EM contraction rate estimation and spurious fixed-point detection.

Computational Budget. Table 20 summarises the computational cost per daily recalibration.

Table 20: Computational budget (T=5,000, K=5, p=30).

Component	**Time**
HSIC-Lasso (500 $\rightarrow$ 30 features)	$\sim$8 s
Sticky HDP-HMM (500 iterations)	$\sim$50 s
PC + HSIC per regime ($\times$5)	$\sim$75 s
SCIT e-value updates	$<$1 s
Shield synthesis (K^2=25 vertices)	$\sim$12.5 min
State abstraction verification	$<$1 s
Mean-variance optimisation	$<$1 s
Total (daily recalibration)	$\sim$15 min
Per-action shield check	$<$1 ms

M Walk-Forward Backtest Details

M.1 Baseline Strategy Descriptions

The following strategies are compared on the same synthetic data streams, all operating on discrete action levels $\{-3, \ldots, 3\}$:

- **Buy-and-hold (B&H)**: Constant position at action level +1.
- **Momentum**: Time-series momentum with 20-day lookback; action magnitude proportional to trailing return.
- **Mean-reversion (MR)**: Fade moves beyond 1.5σ from 50-day rolling mean.
- **Unshielded MV (UMV)**: Mean-variance optimiser with 60-day lookback, no safety shield.
- **Oracle**: Uses ground-truth regime labels with predefined per-regime actions.
- **CSAT**: Full pipeline (Sticky HDP-HMM + PC-HSIC + SCIT + posterior-predictive shield with drawdown-recovery specification).

M.2 Scenario Configurations

All scenarios use T=2000 timesteps, 100-step warmup, and 5 bps per-trade transaction costs. The crisis scenario injects a flash crash at t=1000 with magnitude -15% over 10 steps followed by 50 steps of elevated volatility.

Table 21: Walk-forward backtest scenario configurations.

Scenario	**Asset**	K	p	ν	**GARCH**
Equity	Equity	3	8	5.0	Yes
FX	FX	2	6	8.0	Yes
Crypto	Crypto	4	5	3.0	Yes
Crisis	Equity	3	8	4.0	Yes
Low SNR	Equity	3	10	6.0	Yes

Table 22: Max drawdown across scenarios and strategies.

	B&H	Mom.	MR	UMV	Oracle	CSAT
Equity	22.8%	17.0%	16.4%	26.3%	43.5%	49.5%
FX	20.3%	43.8%	11.9%	37.8%	30.1%	43.4%
Crypto	14.1%	36.9%	7.1%	41.6%	21.8%	43.8%
Crisis	29.5%	29.4%	18.4%	35.9%	30.7%	48.2%
Low SNR	19.1%	42.1%	5.3%	45.2%	8.2%	31.0%

M.3 Extended Results

CSAT's higher max drawdowns reflect the shield permitting aggressive positions in certain states. This is an inherent tradeoff: the shield guarantees *eventual* recovery from drawdowns (bounded liveness) but does not cap instantaneous drawdown magnitude. Combining the shield with position-sizing constraints would address this.

References

[1] M. Alshiekh, R. Bloem, R. Ehlers, B. Könighofer, S. Niekum, and U. Topcu. Safe reinforcement learning via shielding. In *AAAI Conference on Artificial Intelligence*, 2018.

[2] E. Bareinboim and J. Pearl. Causal inference and the data-fusion problem. *Proceedings of the National Academy of Sciences*, 113(27):7345–7352, 2016.

[3] O. Catoni. *PAC-Bayesian Supervised Classification: The Thermodynamics of Statistical Learning*. IMS Lecture Notes, 2007.

[4] R. Christiansen and J. Peters. Switching regression models and causal inference in the presence of discrete latent variables. *JMLR*, 21(41):1–46, 2020.

[5] E. Clarke, O. Grumberg, S. Jha, Y. Lu, and H. Veith. Counterexample-guided abstraction refinement. In *CAV*, pages 154–169, 2000.

[6] M. Cubuktepe, N. Jansen, S. Junges, J.-P. Katoen, and U. Topcu. Scenario-based verification of uncertain MDPs. In *TACAS*, 2021.

[7] A. P. Dempster, N. M. Laird, and D. B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *JRSS-B*, 39(1):1–38, 1977.

[8] E. B. Fox, E. B. Sudderth, M. I. Jordan, and A. S. Willsky. A sticky HDP-HMM with application to speaker diarization. *Annals of Applied Statistics*, 5(2A):1020–1056, 2011.

[9] J. Grand-Clément and M. Petrik. Convex formulations for non-rectangular robust Markov decision processes. *Operations Research*, 2022.

[10] A. Gretton, K. Fukumizu, C. H. Teo, L. Song, B. Schölkopf, and A. Smola. A kernel statistical test of independence. In *NeurIPS*, 2007.

[11] P. Grünwald, R. de Heide, and W. Koolen. Safe testing. *JRSS-B*, 86(5):1091–1128, 2024.

[12] J. D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989.

[13] G. N. Iyengar. Robust dynamic programming. *Mathematics of Operations Research*, 30(2):257–280, 2005.

[14] N. Jansen, B. Könighofer, S. Junges, A. Serban, and R. Bloem. Safe reinforcement learning using probabilistic shields. In *CONCUR*, 2020.

[15] K.-S. Jun and F. Orabona. Parameter-free online convex optimization with sub-exponential noise. In *COLT*, 2019.

[16] D. A. McAllester. PAC-Bayesian model averaging. In *COLT*, 1999.

[17] A. Nilim and L. El Ghaoui. Robust control of Markov decision processes with uncertain transition matrices. *Operations Research*, 53(5):780–798, 2005.

[18] J. Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2nd edition, 2009.

[19] J. Peters, P. Bühlmann, and N. Meinshausen. Causal inference by using invariant prediction: identification and confidence intervals. *JRSS-B*, 78(5):947–1012, 2016.

[20] N. Pfister, P. Bühlmann, and J. Peters. Invariant causal prediction for sequential data. *JASA*, 114(527):1264–1276, 2019.

[21] P. Spirtes, C. Glymour, and R. Scheines. *Causation, Prediction, and Search*. MIT Press, 2nd edition, 2000.

[22] J. Van Gael, Y. Saatci, Y. W. Teh, and Z. Ghahramani. Beam sampling for the infinite hidden Markov model. In *ICML*, 2008.

[23] M. Yamada, W. Jitkrittum, L. Sigal, E. P. Xing, and M. Sugiyama. High-dimensional feature selection by feature-wise kernelized Lasso. *Neural Computation*, 26(1):185–207, 2014.

[24] K. Zhang, J. Peters, D. Janzing, and B. Schölkopf. Kernel-based conditional independence test and application in causal discovery. In *UAI*, 2012.