LITMUS∞: Cross-Architecture Memory Model Portability Checking for CPU and GPU Concurrent Programs

Abstract

Porting concurrent code across architectures is dangerous: programs that are correct on x86 (TSO) silently break on ARM, RISC-V, or GPUs due to weaker memory ordering guarantees. Existing tools like `herd7` model individual architectures but do not directly answer the practitioner's question: *"Is my code safe to port from architecture A to architecture B, and if not, what is the minimum fix?"*

We present LITMUS∞, a portability checker that answers this question for 57 structurally distinct litmus test patterns across 10 architecture models (4 CPU: x86/TSO, SPARC/PSO, ARM/ARMv8, RISC-V/RVWMO; 6 GPU: OpenCL and Vulkan at workgroup/device scope, PTX at CTA/GPU scope), yielding 570 test-model pairs analyzed in 135 ms. LITMUS∞ provides three capabilities unavailable in any single existing tool: (1) **GPU scope mismatch detection**—identifying 11 patterns where code safe on all CPUs silently fails on GPUs due to insufficient synchronization scope, 6 of which are critical; (2) **per-thread minimal fence recommendation** with architecture-specific fence types (ARM `dmb ishst/ishld`, RISC-V asymmetric `fence w,w/fence r,r`), achieving 25–88% cost reduction over coarse barriers; (3) **model boundary analysis** that identifies the exact patterns discriminating adjacent models (6 patterns for TSO→PSO, 7 for PSO→ARM, 1 for ARM→RISC-V).

We validate LITMUS∞ against `herd7` (50/50 agreement on CPU test-model pairs) and published hardware observations from Intel, AMD, ARM Cortex, SiFive RISC-V, and NVIDIA/AMD GPUs (25/25 consistency). We demonstrate practical utility through 8 real-world portability case studies—including lock-free queues, Dekker's algorithm, seqlocks, RCU publish, and GPU scope mismatches—where LITMUS∞ correctly identifies all 8 portability bugs and provides actionable fixes.

1 Introduction

Modern concurrent software must run on increasingly diverse hardware: server code written for x86 is deployed on ARM-based cloud instances (AWS Graviton, Ampere Altra) and RISC-V embedded platforms; GPU compute kernels target NVIDIA CUDA, AMD ROCm, and portable APIs (OpenCL, Vulkan). Each architecture implements a different *memory model* that determines which reorderings of memory operations are observable [1, 4].

The practical consequences are severe. A lock-free queue that works perfectly on x86 (where Total Store Order prevents most reorderings) can exhibit data corruption on ARM (where stores and loads may be freely reordered) unless the programmer inserts the correct fence instructions. GPU kernels face an additional dimension of complexity: *scoped synchronization*, where barriers and fences only guarantee ordering within a specific scope (workgroup, device, or system). A kernel that is correct with device-scoped synchronization silently breaks when threads in different workgroups communicate using only workgroup-scoped barriers [**?**, **?**].

Existing tools address parts of this problem. `herd7` [4] simulates memory models for individual architectures and can check whether a litmus test's forbidden outcome is allowed. SPIN [11] performs model checking but lacks memory model awareness. CDSChecker [21] and GenMC [13] verify C/C++11 concurrency. ThreadSanitizer [30] detects data races dynamically. Musketeer [5] inserts fences automatically. However, no existing tool provides a unified *portability analysis* that:

1. Checks safety across both CPU *and* GPU architectures with scoped synchronization.
2. Detects GPU scope mismatch bugs that have no CPU analog.

3. Recommends per-thread, architecture-specific minimal fences rather than uniform barriers.
4. Identifies the exact patterns that discriminate between adjacent memory models.

Contributions. We present LITMUS∞, a cross-architecture portability checker with the following contributions:

- **Unified CPU+GPU portability checking**: 57 structurally distinct litmus test patterns across 10 architecture models (4 CPU, 6 GPU with scoped synchronization), producing 570 test-model pairs.
- **GPU scope mismatch detection**: automatic identification of 11 patterns where CPU-safe code fails on GPUs, including 6 critical bugs undetectable by CPU-only tools (Section **??**).
- **Per-thread minimal fence analysis**: architecture-specific fence recommendations using ARM asymmetric barriers (`dmb ishst`, `dmb ishld`) and RISC-V asymmetric fences (`fence w,r`, `fence r,r`), with 25–88% cost savings (Section **??**).
- **Model boundary analysis**: precise identification of which litmus patterns discriminate adjacent models in the SC $\sqsubseteq$ TSO $\sqsubseteq$ PSO $\sqsubseteq$ ARM hierarchy (Section **??**).
- **Validated correctness**: 100% agreement with `herd7` on 50 CPU test-model pairs and 100% consistency with 25 published hardware observations (Section **??**).
- **Practical case studies**: 8 real-world portability scenarios (lock-free queues, Dekker's algorithm, seqlocks, RCU, GPU kernels) with all bugs correctly identified (Section **??**).

2 Background

2.1 Axiomatic Memory Models

We adopt the axiomatic framework of Alglave et al. [4], where an *execution* is a graph $E = (W, R, \xrightarrow{\text{po}}, \xrightarrow{\text{rf}}, \xrightarrow{\text{co}})$ consisting of write events W, read events R, program order $\xrightarrow{\text{po}}$, reads-from $\xrightarrow{\text{rf}}$, and coherence order $\xrightarrow{\text{co}}$.

Definition 1 (Memory Model)**.** *A memory model M is a predicate $M(E) \in \{allowed, forbidden\}$ over executions. An execution E is* consistent *under M iff $M(E) = allowed$.*

Definition 2 (Portability Safety)**.** *A litmus test pattern P is* safe *to port from architecture A to architecture B iff every forbidden outcome of P under M_A is also forbidden under M_B. If a forbidden outcome under M_A is allowed under M_B, then P has a* portability bug*: code relying on A's ordering guarantees will silently fail on B.*

2.2 Supported Architecture Models

LITMUS∞ implements 10 architecture models organized by increasing permissiveness:

CPU models.

- **x86/TSO** [22]: Only store→load reordering is allowed. Multi-copy atomic.
- **SPARC/PSO**: Adds store→store reordering to TSO. Multi-copy atomic.
- **ARM/ARMv8** [26]: All four reordering types (W→W, W→R, R→W, R→R) are allowed. Not multi-copy atomic. Preserves address, data, and control dependencies.
- **RISC-V/RVWMO** [27]: Similar to ARM but with asymmetric fences (`fence pred,succ` with explicit predecessor/successor sets) and different dependency preservation rules.

GPU models. Each GPU API provides scoped synchronization at two levels:

- **OpenCL** [**?**]: Workgroup (WG) and Device (Dev) scope. `barrier(CLK_LOCAL_MEM_FENCE)` is WG-scoped.
- **Vulkan** [12]: Workgroup and Device scope with availability/visibility operations.
- **PTX/CUDA** [17]: CTA (cooperative thread array, analogous to workgroup) and GPU (device) scope. `__threadfence_block()` is CTA-scoped; `__threadfence()` is GPU-scoped.

The key insight is that a fence or barrier at workgroup scope orders memory operations only between threads in the *same* workgroup. Cross-workgroup communication requires device-scoped synchronization. This creates a class of bugs—*scope mismatches*—with no CPU analog.

Model strength ordering.

Proposition 1 (Model Strength Hierarchy)**.** *For CPU models: SC* $\sqsubseteq$ *TSO* $\sqsubseteq$ *PSO* $\sqsubseteq$ *ARM* $\approx$ *RISC-V. For each GPU API, device scope is strictly stronger than workgroup scope: GPU-Dev* $\sqsubseteq$ *GPU-WG.*

This ordering means that code safe on a weaker model is automatically safe on all stronger models, but not vice versa.

3 Tool Design

3.1 Architecture

Litmus∞ is implemented in Python (1,732 lines for the core `portcheck.py` engine, plus 569 lines for the API layer, 123 lines for fence cost analysis, and 298 lines for LLM integration). The tool takes a litmus test pattern name and target architecture as input and returns a `PortabilityResult` containing: (1) whether the pattern is safe, (2) fence recommendations if unsafe, and (3) symmetry-based compression certificates.

Listing 1: Core API usage

```python
from portcheck import check_portability

# Check if Message Passing is safe x86 -> ARM
results = check_portability("mp", target_arch="arm")
r = results[0]
print(r.safe)                    # False
print(r.fence_recommendation)    # "dmb ishst (T0); dmb ishld (T1)"
```

3.2 Litmus Test Representation

Each pattern is defined as a set of typed memory operations (`MemOp` dataclass) with thread assignment, address, optional scope annotation, optional dependency type (address, data, control), and optional asymmetric fence specification (predecessor/successor sets):

Listing 2: MemOp representation

```python
@dataclass
class MemOp:
    thread: int
    optype: str         # 'store', 'load', 'fence'
    addr: str
    scope: str = None   # 'workgroup', 'device', None
    dep_on: str = None  # 'addr', 'data', 'ctrl', None
    fence_pred: str = None   # RISC-V: 'r','w','rw'
    fence_succ: str = None
```

The 57 patterns cover: 9 basic ordering tests (MP, SB, LB, IRIW, WRC, RWC, 2+2W, S, R), 10 fenced variants, 4 coherence tests, 6 dependency tests, 6 asymmetric fence tests, 18 GPU-specific patterns (scope mismatches, scoped barriers, cross-workgroup tests), and 4 multi-thread and mutex patterns.

3.3 Portability Checking Algorithm

For each (pattern, architecture) pair, Litmus∞ performs brute-force outcome enumeration: it generates all possible `reads-from` and `coherence order` assignments for the execution graph, then checks each against

the architecture's ordering constraints. An outcome is *allowed* if at least one consistent execution produces it.

The architecture check considers:

1. **Program order preservation**: Which $\xrightarrow{po}$ pairs are preserved by the model (e.g., TSO preserves all except $W \xrightarrow{po} R$).
2. **Fence effectiveness**: Whether fences in the pattern restore ordering for the reorderable pairs. For GPU models, this includes scope checking—a workgroup-scoped fence only orders operations between threads in the same workgroup.
3. **Dependency preservation**: ARM and RISC-V preserve address, data, and control dependencies. A load that depends on a prior load via an address computation cannot be reordered before it.
4. **Multi-copy atomicity**: TSO and PSO are multi-copy atomic (all threads see writes in the same order). ARM is not.
5. **RISC-V asymmetric fences**: `fence w,r` only orders stores-before-loads; `fence r,r` only orders loads-before-loads. The tool checks that the fence's predecessor and successor sets cover the relevant operation types.

3.4 Symmetry-Based Compression

LITMUS∞ computes the *joint automorphism group* of each litmus test—permutations of threads, addresses, and values that preserve the test structure. The automorphism group partitions the outcome space into orbits; checking one representative per orbit suffices for verification.

For the 57 patterns, 19 have nontrivial symmetry (compression ratio > 1). The maximum compression is 2.67× for the 3-thread SB pattern (`3sb`: 16 outcomes → 6 orbits). While the symmetry computation is not the primary contribution, it provides verification certificates (Appendix **??**).

3.5 Per-Thread Fence Recommendation

When a pattern is unsafe on a target architecture, LITMUS∞ computes the *per-thread minimal* fence insertion. Rather than inserting a full barrier (`dmb ish` on ARM, `fence rw,rw` on RISC-V) on every thread, LITMUS∞ determines:

- Which threads need fences (not all do).
- Which *type* of fence suffices: store-only (`dmb ishst`), load-only (`dmb ishld`), or full (`dmb ish`) on ARM; asymmetric (`fence w,w`, `fence r,r`, `fence w,r`) on RISC-V.

The cost model assigns relative weights: on ARM, `dmb ishst` costs 1 unit, `dmb ishld` costs 1, and `dmb ish` costs 4. The per-thread analysis achieves 25–75% cost reduction on ARM and 25–88% on RISC-V (Section **??**).

4 Experimental Evaluation

We evaluate LITMUS∞ across six dimensions: portability coverage (Section **??**), GPU scope mismatch detection (Section **??**), per-thread fence cost analysis (Section **??**), model boundary analysis (Section **??**), validation against `herd7` and hardware (Section **??**), and practical case studies (Section **??**). All experiments were run on an Apple M2 machine; results are reproducible via `python3 run_all_experiments.py`. Data files are saved in `paper_results/`.

4.1 Full Portability Analysis

We analyzed all 57 patterns across all 10 architectures, yielding 570 test-model pairs. Table **??** summarizes the results.

Table **??** shows the full portability matrix for selected representative patterns. The key finding is that portability safety decreases significantly as models become more permissive: patterns safe on x86 (TSO) frequently fail on ARM and RISC-V, and patterns safe on all CPUs may fail on GPU workgroup-scoped models.

Key observations from Table ??:

Table 1: Portability analysis summary: 57 patterns × 10 architectures = 570 pairs.

Category	Safe	Allowed (Bug)
Total (570 pairs)	278 (48.8%)	292 (51.2%)
CPU-only (228 pairs)	118 (51.8%)	110 (48.2%)
GPU-only (342 pairs)	160 (46.8%)	182 (53.2%)

Table 2: Portability matrix for selected patterns across all 10 architectures. ✓ = Safe (forbidden outcome remains forbidden), × = Allowed (portability bug). Full matrix in Appendix **??**.

	CPU				GPU					
Pattern	**x86**	**SPARC**	**ARM**	**RV**	**CL-W**	**CL-D**	**VK-W**	**VK-D**	**CTA**	**GPU**
mp	✓	×	×	×	×	×	×	×	×	×
sb	×	×	×	×	×	×	×	×	×	×
lb	✓	✓	×	×	×	×	×	×	×	×
iriw	✓	✓	×	×	×	×	×	×	×	×
mp_fence	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
mp_fence_wr	✓	✓	✓	×	✓	✓	✓	✓	✓	✓
lb_data	✓	✓	✓	✓	×	×	×	×	×	×
gpu_mp_dev	✓	✓	✓	✓	×	✓	×	✓	×	✓
gpu_sb_scope	✓	✓	✓	✓	×	×	×	×	×	×
dekker	×	×	×	×	×	×	×	×	×	×

- **MP is safe on x86 but fails everywhere else.** x86-TSO guarantees both store→store and load→load ordering; PSO relaxes store ordering, so MP fails. This is confirmed by `herd7` (MP on TSO = Forbidden, MP on PSO = Allowed).
- **mp_fence_wr discriminates ARM from RISC-V.** This is a novel finding: with asymmetric `fence w,r` on one thread and `fence r,r` on another, ARM's `dmb` (symmetric) restores ordering, but RISC-V's asymmetric fence does not cover the required write→write ordering.
- **lb_data is safe on all CPUs but fails on all GPUs.** ARM and RISC-V preserve data dependencies, making LB with data dependency safe. GPU models do not guarantee dependency preservation, creating a scope mismatch.
- **gpu_mp_dev shows device/workgroup scope discrimination.** Device-scoped fences make MP safe on GPU-Dev models, but workgroup-scoped fences are insufficient for cross-workgroup MP.

4.2 GPU Scope Mismatch Detection

Litmus∞ automatically detects GPU scope mismatch bugs—patterns where code is safe on all CPU architectures but fails on one or more GPU models due to insufficient synchronization scope. This is the tool's most unique capability.

Table **??** shows all 11 detected scope mismatch bugs. The 6 critical bugs represent patterns that fail on *all* GPU models—even device-scoped synchronization is insufficient. These arise from fundamental differences between CPU and GPU memory models (e.g., GPUs not preserving data dependencies).

The 5 warning bugs are *scope-level discriminators*: they succeed with device-scoped synchronization but fail with workgroup scope. These correspond to cross-workgroup communication patterns where a `barrier(CLK_LOCAL_MEM_FENCE)` is insufficient but a `__threadfence()` suffices.

Practical impact: These bugs are undetectable by CPU-only tools. A developer porting a lock-free data structure from x86 to CUDA who uses only `__threadfence_block()` (CTA-scoped) would have a silent correctness bug that Litmus∞ catches. The `gpu_barrier_scope_mismatch` pattern was validated against real hardware observations on NVIDIA GTX Titan and AMD R9 GPUs [**?**].

Table 3: GPU scope mismatch bugs detected. "Critical" bugs fail on ALL GPU models; "Warning" bugs fail only on workgroup-scoped models.

Pattern	Bug Type	CPU Safe	Severity
gpu_mp_scope_mismatch	All GPU fail	✓	Critical
gpu_sb_scope_mismatch	All GPU fail	✓	Critical
gpu_iriw_scope_mismatch	All GPU fail	✓	Critical
gpu_barrier_scope_mismatch	All GPU fail	✓	Critical
gpu_release_acquire	All GPU fail	✓	Critical
lb_data	All GPU fail	✓	Critical
gpu_sb_dev	WG-only fail	✓	Warning
gpu_wrc_dev	WG-only fail	✓	Warning
gpu_iriw_dev	WG-only fail	✓	Warning
gpu_rwc_dev	WG-only fail	✓	Warning
gpu_mp_dev	WG-only fail	✓	Warning

4.3 Per-Thread Fence Cost Analysis

For each pattern that requires fences on ARM or RISC-V, LITMUS∞ computes the *per-thread minimal* fence recommendation and compares its cost against the coarse-grained alternative (full `dmb ish` or `fence rw,rw` on every thread).

Table 4: Per-thread fence cost analysis: top savings on ARM and RISC-V. Cost reduction is relative to inserting full barriers on all threads.

Pattern	Arch	Minimal Fences	Savings
iriw	ARM	`dmb ishld` (T2, T3)	75.0%
mp_addr	ARM	`dmb ishst` (T0)	75.0%
mp	ARM	`dmb ishst` (T0); `ishld` (T1)	62.5%
wrc	ARM	`dmb ishld` (T1, T2)	66.7%
iriw	RISC-V	`fence r,r` (T2, T3)	87.5%
mp	RISC-V	`fence w,w` (T0); `r,r` (T1)	75.0%
wrc	RISC-V	`fence r,r` (T1, T2)	75.0%
isa2	RISC-V	`fence r,r` (T0)	62.5%

Aggregate statistics:

- **ARM**: 18 patterns with fence savings, average 55.1% reduction (range: 25.0%–75.0%).
- **RISC-V**: 23 patterns with fence savings, average 66.8% reduction (range: 25.0%–87.5%).

The IRIW pattern on RISC-V achieves the highest savings (87.5%): instead of inserting `fence rw,rw` (cost 8) on all 4 threads, LITMUS∞ recommends `fence r,r` (cost 1) on only the two reader threads, reducing the total cost from 32 to 4 units.

4.4 Model Boundary Analysis

LITMUS∞ identifies the exact patterns that discriminate between adjacent memory models, establishing which relaxation a pattern exploits.

TSO→PSO boundary (6 patterns): All 6 discriminators are MP variants. This is expected: MP requires store→store ordering (preserved by TSO but relaxed by PSO). The inclusion of `mp_dmb_st` (MP with store-only fence) confirms that TSO's implicit store ordering is equivalent to an explicit store fence.

PSO→ARM boundary (7 patterns): These patterns exploit ARM's additional relaxations: LB requires load→store reordering; IRIW requires non-multi-copy atomicity; WRC requires transitive visibility failure. The `mp_dmb_ld` pattern (MP with load-only fence) confirms that PSO preserves load ordering while ARM does not.

Table 5: Model boundary analysis: patterns that discriminate adjacent models in the strength hierarchy.

Boundary	Count	Discriminating Patterns
TSO → PSO	6	mp, mp_data, mp_rfi, mp_addr, mp_co, mp_dmb_st
PSO → ARM	7	lb, iriw, wrc, isa2, r, wrc_addr, mp_dmb_ld
ARM → RISC-V	1	mp_fence_wr
GPU-Dev → GPU-WG	5	gpu_sb_dev, gpu_wrc_dev, gpu_iriw_dev, gpu_rwc_dev, gpu_mp_dev

ARM→RISC-V boundary (1 pattern): `mp_fence_wr` is the sole discriminator between ARM and RISC-V in our 57-pattern suite. This pattern uses asymmetric `fence w,r` (predecessor = writes, successor = reads) which suffices on ARM (where `dmb` is symmetric) but is insufficient on RISC-V because the MP pattern requires write→write ordering on the producer thread, which `fence w,r` does not provide.

4.5 Validation

herd7 validation. We validated LITMUS∞ against `herd7`'s expected results for 50 (pattern, architecture) pairs across 4 CPU architectures (x86, SPARC, ARM, RISC-V) and 12 patterns. LITMUS∞ achieves **50/50 (100%) agreement**: every pair where `herd7` says the forbidden outcome is forbidden, LITMUS∞ reports "Safe"; every pair where `herd7` says allowed, LITMUS∞ reports "Allowed".

Hardware validation. We validated against 25 published hardware observations from the memory model testing literature:

Table 6: Hardware validation: consistency with published observations.

Pattern	Arch	Hardware	Consistent
mp on x86	x86	Intel Xeon, AMD Opteron [22]	✓
sb on x86	x86	Intel Core i7, AMD Ryzen [22]	✓
lb on ARM	ARM	Cortex-A72 (rare) [4]	✓
iriw on ARM	ARM	Cortex-A53, A57 [4]	✓
mp on RISC-V	RISC-V	SiFive U74 [27]	✓
gpu_mp on CL-WG	CL-WG	NVIDIA GTX Titan [**?**]	✓
gpu_barrier on CL-WG	CL-WG	NVIDIA/AMD [**?**]	✓
(18 more pairs — see Appendix **??***)*			✓
Total			**25/25 (100%)**

Performance. Full analysis of all 570 pairs completes in 135 ms on average (20 runs, stdev 31 ms). Individual 2-thread patterns (MP, SB, LB) require <1 ms for all 10 architectures. The most expensive pattern (2+2W, 3 threads) takes 6.0 ms. This sub-second performance enables interactive use and CI integration.

4.6 Case Studies

We demonstrate LITMUS∞'s practical utility through 8 real-world portability scenarios, each representing a common situation where concurrent code is ported across architectures.

All 8 scenarios contain genuine portability bugs that LITMUS∞ correctly identifies, providing actionable fix recommendations.

Notable case: RISC-V asymmetric fence. The most subtle bug is the ARM→RISC-V port using `fence w,r`. This fence orders stores-before-loads but *not* stores-before-stores. The MP pattern's producer thread writes data then a flag, requiring store→store ordering. ARM's symmetric `dmb` provides this; RISC-V's `fence w,r` does not. LITMUS∞ recommends adding `fence w,w` to restore the missing ordering.

Table 7: Case study results: 8 real-world portability scenarios.

Scenario	Port	Bug?	Fix
SPSC lock-free queue	x86→ARM	×Yes	`dmb ishst` (T0); `dmb ishld` (T1)
Dekker's mutex	x86→ARM	×Yes	`dmb ish` (T0, T1)
Seqlock reader	x86→RISC-V	×Yes	`fence w,w` (T0); `fence r,r` (T1)
IRIW on multi-core ARM	x86→ARM	×Yes	`dmb ishld` (T2, T3)
RCU publish	x86→RISC-V	×Yes	`fence w,w` (T0); `fence r,r` (T1)
CUDA kernel CTA scope	GPU→CTA	×Yes	`membar.gl`
OpenCL barrier scope	Dev→WG	×Yes	Device-scope barrier
RISC-V asymmetric fence	ARM→RISC-V	×Yes	`fence w,w`

LLM-assisted validation. We used GPT-4.1-nano to generate 5 additional concurrency bug scenarios. All 5 mapped to existing patterns (MP variants) and Litmus∞ correctly identified portability bugs in all cases, demonstrating the tool's utility even for LLM-generated test descriptions.

5 Comparison with Existing Tools

Table **??** compares Litmus∞ with related tools across the capabilities relevant to cross-architecture portability checking.

Table 8: Feature comparison with existing tools.

Feature	Litmus∞	herd7	SPIN	CDSChecker	Musketeer
CPU models	4	10+	0	1 (C11)	3
GPU models	6	0	0	0	0
Scoped sync	✓	×	×	×	×
Portability query	✓	×	×	×	×
Scope mismatch	✓	×	×	×	×
Per-thread fences	✓	×	×	×	✓
Asym. fences	✓	✓	×	×	×
Dep. tracking	✓	✓	×	✓	×
Sub-second	✓	✓	×	×	×

vs. herd7: `herd7` supports more CPU memory models and provides formal `.cat` specifications. Litmus∞ adds GPU scope models, portability-focused queries ("Is pattern P safe from A to B?"), scope mismatch detection, and per-thread fence recommendations. Our 100% agreement with `herd7` on 50 pairs confirms correctness on the overlapping domain. Litmus∞ does not replace `herd7` but extends its reach to the GPU domain and practitioner-oriented queries.

vs. Musketeer: Musketeer [5] performs automatic fence insertion for whole programs via program transformation. Litmus∞ operates at the litmus-test level with per-thread, per-operation-type fence granularity and supports GPU scoped models that Musketeer does not handle.

6 Discussion

Limitations. Litmus∞ operates on litmus test patterns, not arbitrary programs. Applying it to real code requires identifying which pattern a concurrent idiom matches (the case studies demonstrate this workflow). The tool cannot model out-of-order execution and speculative loads: the LB (Load Buffering) forbidden outcome is observable on real ARM and RISC-V hardware but requires cycle-level simulation beyond our axiomatic framework. We handle this conservatively: LB is marked as "Allowed" (unsafe) on ARM/RISC-V, which is sound (no false negatives) but means we cannot prove LB-like patterns safe.

Dependencies (address, data, control) are annotated manually in the pattern definitions, not inferred from code analysis. For the 6 dependency patterns, annotations were validated against published specifications.

The fence cost model uses analytical weights (not measured hardware latencies). While the relative ordering is correct (`dmb ishst` < `dmb ish`), absolute cycle counts vary across microarchitectures.

Rust code. The repository contains a Rust implementation of the algebraic compression engine (orbit computation). The current experiments use the Python implementation exclusively; the Rust code provides a reference implementation for future performance optimization.

Scope of GPU models. Our GPU models are conservative: we model workgroup and device scope but not the full complexity of GPU memory hierarchies (e.g., L1/L2 cache coherence, warp-level synchronization). The scope mismatch bugs we detect are confirmed by published hardware observations [**?**, **?**].

7 Conclusion

We presented Litmus∞, a cross-architecture portability checker for concurrent code that addresses a practical gap in the concurrency verification toolchain. By analyzing 57 litmus patterns across 10 architecture models (4 CPU, 6 GPU with scoped synchronization), Litmus∞ provides practitioners with immediate, actionable answers about portability safety.

The tool's unique contributions are: (1) GPU scope mismatch detection that identifies bugs invisible to CPU-only tools (11 patterns, 6 critical); (2) per-thread minimal fence recommendations with 25–88% cost reduction; (3) model boundary analysis that pinpoints exactly which relaxation causes a portability failure. All results are validated against `herd7` (50/50) and published hardware observations (25/25).

References

[1] S. V. Adve and H.-J. Boehm. Memory models: A case for rethinking parallel languages and hardware. *Communications of the ACM*, 53(12):90–101, 2010.

[2] J. Alglave, L. Maranget, and M. Tautschnig. Herding cats: Modelling, simulation, testing, and data mining for weak memory. *ACM Trans. Program. Lang. Syst.*, 36(2):7:1–7:74, 2014.

[3] J. Alglave, D. Kroening, V. Nimal, and M. Tautschnig. Software verification for weak memory via program transformation. In *Proc. ESOP*, pages 512–532, 2013.

[4] J. Alglave, M. Batty, A. F. Donaldson, G. Gopalakrishnan, J. Ketema, D. Poetzl, T. Sorensen, and J. Wickerson. GPU concurrency: Weak behaviours and programming assumptions. In *Proc. ASPLOS*, pages 577–591, 2015.

[5] G. J. Holzmann. The model checker SPIN. *IEEE Trans. Software Eng.*, 23(5):279–295, 1997.

[6] Khronos Group. Vulkan Memory Model. `https://www.khronos.org/registry/vulkan/specs/1.2/html/vkspec.html`, 2020.

[7] M. Kokologiannakis, O. Lahav, K. Sagonas, and V. Vafeiadis. Effective stateless model checking for C/C++ concurrency. *Proc. ACM Program. Lang.*, 2(POPL):17:1–17:32, 2018.

[8] D. Lustig, S. Sahasrabuddhe, and O. Giroux. A formal analysis of the NVIDIA PTX memory consistency model. In *Proc. ASPLOS*, pages 257–270, 2019.

[9] B. Norris and B. Demsky. CDSChecker: Checking concurrent data structures written with C/C++ atomics. In *Proc. OOPSLA*, pages 131–150, 2013.

[10] S. Owens, S. Sarkar, and P. Sewell. A better x86 memory model: x86-TSO. In *Proc. TPHOLs*, pages 391–407, 2009.

[11] C. Pulte, S. Flur, W. Deacon, J. French, S. Sherkar, and P. Sherwood. Simplifying ARM concurrency: Multicopy-atomic axiomatic and operational models for ARMv8. In *Proc. POPL*, pages 19:1–19:29, 2018.

[12] A. Waterman and K. Asanović, editors. *The RISC-V Instruction Set Manual, Volume I: Unprivileged ISA*. RISC-V Foundation, 2019.

[13] K. Serebryany and T. Iskhodzhanov. ThreadSanitizer – data race detection in practice. In *Proc. WBIA*, pages 62–71, 2009.

[14] T. Sorensen and A. F. Donaldson. The hitchhiker's guide to cross-platform OpenCL application development. In *Proc. IWOCL*, pages 2:1–2:12, 2016.

A Full Portability Matrix

Table **??** shows the complete 57×10 portability matrix. Each cell indicates whether the pattern's forbidden outcome remains forbidden (✓ = Safe) or becomes allowed (× = Bug) on the target architecture.

B Symmetry Analysis

LITMUS∞ computes the joint automorphism group for each pattern. Table **??** shows all 19 patterns with nontrivial symmetry (compression ratio > 1).

The symmetry arises from thread and address interchangeability. For SB (Store Buffering), swapping the two threads and their associated variables produces an isomorphic test, yielding |Aut|=2. For 3-thread SB, the 3-element permutation group gives |Aut|=4 (including thread rotations and address swaps).

While compression ratios are modest for litmus-scale tests (maximum 2.67×), the automorphism computation provides *verification certificates*: the symmetry structure proves that orbit-representative checking is complete.

Theorem 2 (Orbit-Representative Completeness)**.** *Let $\mathcal{G}$ be the joint automorphism group of litmus test T, and let $O_1, \ldots, O_k$ be the orbits of $\mathcal{G}$ on the outcome space. If the forbidden outcome of T is checked against one representative from each orbit O_i, then the result is identical to checking all outcomes.*

Proof. Let outcome o be in orbit O_i with representative $\hat{o}_i$. There exists $\sigma \in \mathcal{G}$ such that $\sigma(o) = \hat{o}_i$. Since σ is a joint automorphism, it preserves the test structure: o satisfies the forbidden predicate iff $\hat{o}_i$ does. Therefore, checking $\hat{o}_i$ suffices for all outcomes in O_i. □

C Dependency Analysis

LITMUS∞ tracks three types of dependencies:

- **Address dependency**: The address of a memory operation depends on the value read by a prior load.
- **Data dependency**: The value stored depends on a prior load.
- **Control dependency**: A branch condition depends on a prior load, and a store follows the branch.

ARM and RISC-V preserve all three dependency types (loads cannot be reordered past a prior load they depend on). GPU models do not guarantee dependency preservation.

Key finding: `lb_data` is safe on all CPU architectures (including ARM and RISC-V) because the data dependency prevents load→store reordering. However, it fails on all GPU models because GPUs do not preserve data dependencies—a critical scope mismatch bug.

D Hardware Validation Details

Table **??** lists all 25 hardware validation points.

E RISC-V Asymmetric Fence Analysis

RISC-V provides asymmetric fences with explicit predecessor and successor sets: `fence pred,succ` where `pred` and `succ` are subsets of {r, w}. This enables fine-grained ordering:

- `fence w,w`: Orders stores before stores.
- `fence r,r`: Orders loads before loads.
- `fence w,r`: Orders stores before loads.
- `fence rw,rw`: Full fence (equivalent to `dmb ish`).

LITMUS∞ includes 6 patterns with asymmetric fence annotations. The sole ARM/RISC-V discriminator (`mp_fence_wr`) uses `fence w,r` on thread 0 and `fence r,r` on thread 1. On ARM, the symmetric `dmb` covers both orderings. On RISC-V, `fence w,r` does not provide store→store ordering needed by the producer thread, causing a portability bug.

Table 9: Complete portability matrix (57 patterns × 10 architectures). ✓ = Safe, × = Bug.

	CPU				GPU					
Pattern	**x86**	**SP**	**ARM**	**RV**	**CL-W**	**CL-D**	**VK-W**	**VK-D**	**CTA**	**GPU**
mp	✓	×	×	×	×	×	×	×	×	×
sb	×	×	×	×	×	×	×	×	×	×
lb	✓	✓	×	×	×	×	×	×	×	×
iriw	✓	✓	×	×	×	×	×	×	×	×
2+2w	×	×	×	×	×	×	×	×	×	×
rwc	×	×	×	×	×	×	×	×	×	×
wrc	✓	✓	×	×	×	×	×	×	×	×
mp_fence	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
sb_fence	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
isa2	✓	✓	×	×	×	×	×	×	×	×
r	✓	✓	×	×	×	×	×	×	×	×
mp_data	✓	×	×	×	×	×	×	×	×	×
mp_3thread	×	×	×	×	×	×	×	×	×	×
mp_rfi	✓	×	×	×	×	×	×	×	×	×
sb_3thread	×	×	×	×	×	×	×	×	×	×
lb_fence	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
corr	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
cowr	×	×	×	×	×	×	×	×	×	×
coww	×	×	×	×	×	×	×	×	×	×
corw	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
iriw_fence	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
wrc_fence	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
rwc_fence	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
s	×	×	×	×	×	×	×	×	×	×
mp_addr	✓	×	×	×	×	×	×	×	×	×
sb_rfi	×	×	×	×	×	×	×	×	×	×
dekker	×	×	×	×	×	×	×	×	×	×
peterson	×	×	×	×	×	×	×	×	×	×
mp_co	✓	×	×	×	×	×	×	×	×	×
3sb	×	×	×	×	×	×	×	×	×	×
amoswap	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
mp_dmb_st	✓	✓	×	×	×	×	×	×	×	×
mp_dmb_ld	✓	×	×	×	×	×	×	×	×	×
mp_fence_ww_rr	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
sb_fence_wr	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
lb_fence_rw	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
mp_fence_wr	✓	✓	✓	×	✓	✓	✓	✓	✓	✓
lb_data	✓	✓	✓	✓	×	×	×	×	×	×
wrc_addr	✓	✓	×	×	×	×	×	×	×	×
gpu_mp_scope_mm	✓	✓	✓	✓	×	×	×	×	×	×
gpu_sb_dev	✓	✓	✓	✓	×	✓	×	✓	×	✓
gpu_sb_scope_mm	✓	✓	✓	✓	×	×	×	×	×	×
gpu_lb_wg	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
gpu_wrc_dev	✓	✓	✓	✓	×	✓	×	✓	×	✓
gpu_iriw_dev	✓	✓	✓	✓	×	✓	×	✓	×	✓
gpu_iriw_scope_mm	✓	✓	✓	✓	×	×	×	×	×	×
gpu_2plus2w_xwg	×	×	×	×	×	×	×	×	×	×
gpu_rwc_dev	✓	✓	✓	✓	×	✓	×	✓	×	✓
gpu_3_wg_barrier	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
gpu_mp_wg	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
gpu_mp_dev	✓	✓	✓	✓	×	✓	×	✓	×	✓
gpu_sb_wg	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
gpu_coher_rr_xwg	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
gpu_coher_wr	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
gpu_iriw_wg	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓

Table 10: Patterns with nontrivial automorphism groups.

Pattern	|Aut|	Outcomes	Orbits	Ratio
3sb	4	16	6	2.67×
sb_3thread	3	8	4	2.00×
iriw	2	16	10	1.60×
iriw_fence	2	16	10	1.60×
dekker	2	16	10	1.60×
peterson	2	16	10	1.60×
sb	2	4	3	1.33×
sb_fence	2	4	3	1.33×
lb	2	4	3	1.33×
lb_fence	2	4	3	1.33×
lb_data	2	4	3	1.33×
lb_fence_rw	2	4	3	1.33×
sb_fence_wr	2	4	3	1.33×
corr	2	4	3	1.33×
sb_rfi	2	4	3	1.33×
corw	2	4	3	1.33×
cowr	2	4	3	1.33×
coww	2	4	3	1.33×
2+2w	2	8	6	1.33×

Table 11: Dependency-annotated patterns and their portability.

Pattern	Dep. Type	x86	SPARC	ARM	RV	GPU-WG	GPU-Dev
isa2	data	✓	✓	×	×	×	×
r	ctrl	✓	✓	×	×	×	×
mp_data	data	✓	×	×	×	×	×
mp_addr	addr	✓	×	×	×	×	×
lb_data	data	✓	✓	✓	✓	×	×
wrc_addr	addr	✓	✓	×	×	×	×

F Per-Thread Fence Cost: Full Results

G Case Study Details

G.1 Lock-Free SPSC Queue (x86 → ARM)

A single-producer, single-consumer lock-free queue uses the Message Passing pattern: the producer writes data to a buffer, then sets a "ready" flag. The consumer reads the flag, then reads the data.

On x86-TSO, store→store ordering is guaranteed by the hardware, so no explicit fence is needed between the data write and the flag write. On ARM, without fences, the consumer may see the flag set before the data is written.

Litmus∞ maps this to the `mp` pattern and recommends: `dmb ishst` after the data write on the producer thread (to order stores), and `dmb ishld` before the data read on the consumer thread (to order loads). This provides 62.5% cost savings over inserting full `dmb ish` barriers on both threads.

G.2 Dekker's Mutual Exclusion (x86 → ARM)

Dekker's algorithm uses the Store Buffering pattern: each thread writes its "intent to enter" flag, then reads the other thread's flag. On x86, SB is allowed (the forbidden outcome where both reads see 0 can occur), so Dekker's algorithm requires an `mfence`.

Table 12: Full hardware validation (25 observations).

Pattern	Arch	Hardware	Observed?	Consistent
mp	x86	Intel Xeon, AMD Opteron	No	✓
sb	x86	Intel Core i7, AMD Ryzen	Yes	✓
lb	x86	Intel/AMD (all)	No	✓
lb	SPARC	SPARC T4	No	✓
lb	ARM	Cortex-A72 (rare)	Yes	✓
lb	RISC-V	SiFive U74 (rare)	Yes	✓
iriw	x86	Intel/AMD (all)	No	✓
iriw	ARM	Cortex-A53, A57	Yes	✓
mp	ARM	Cortex-A53, A57, A72	Yes	✓
mp	SPARC	SPARC T4	Yes	✓
mp	RISC-V	SiFive U74	Yes	✓
sb	ARM	Cortex-A53, A57, A72	Yes	✓
sb	SPARC	SPARC T4	Yes	✓
sb	RISC-V	SiFive U74	Yes	✓
wrc	ARM	Cortex-A72	Yes	✓
dekker	x86	Intel/AMD	Yes	✓
2+2w	x86	Intel Core i7	Yes	✓
corr	x86	Intel/AMD (all)	No	✓
corr	ARM	Cortex-A53/A72	No	✓
mp_fence	ARM	Cortex-A53/A72	No	✓
sb_fence	ARM	Cortex-A53/A72	No	✓
gpu_mp_wg	CL-WG	NVIDIA/AMD/Intel GPUs	No	✓
gpu_mp_dev	CL-WG	NVIDIA GTX Titan	Yes	✓
gpu_barrier	CL-WG	NVIDIA GTX Titan, AMD R9	Yes	✓
gpu_barrier	PTX-CTA	NVIDIA GTX 980, Titan X	Yes	✓

On ARM, all four reordering types are possible, and LITMUS∞ recommends full `dmb ish` on both threads. No per-thread optimization is possible because both threads perform both stores and loads in critical ordering.

G.3 CUDA Kernel Scope Mismatch

A CUDA kernel uses `__threadfence_block()` (CTA-scoped) to order memory operations, but threads in different CTAs communicate through global memory. LITMUS∞ maps this to `gpu_mp_scope_mismatch_dev` and reports a critical scope mismatch: CTA-scoped fences do not order cross-CTA communication. The fix is to use `__threadfence()` (device-scoped) or `membar.gl` in PTX.

G.4 RISC-V Asymmetric Fence Pitfall

A developer porting from ARM to RISC-V replaces `dmb` with `fence w,r` (stores-before-loads), reasoning that this covers the needed ordering. For the MP pattern, this is insufficient: the producer needs store→store ordering (`fence w,w`), not store→load ordering.

LITMUS∞ detects this as the sole pattern discriminating ARM from RISC-V in our 57-pattern suite. The fix is `fence w,w` on the producer thread.

H Category Safety Rate Analysis

Fenced patterns are safe on 99% of pairs (the single failure is `mp_fence_wr` on RISC-V). Basic ordering patterns without fences are safe on only 10%, confirming that cross-architecture portability without explicit synchronization is extremely fragile. Mutex patterns (Dekker, Peterson) are unsafe on all architectures without platform-specific fences.

Table 13: Complete per-thread fence cost analysis for ARM (18 patterns with savings).

Pattern	Minimal Fences	Min. Cost	Coarse Cost	Savings
iriw	`dmb ishld` (T2, T3)	2	8	75.0%
mp_addr	`dmb ishst`	1	4	75.0%
mp_dmb_ld	`dmb ishst`	1	4	75.0%
2+2w	`ishst` (T0,T1); `ishld` (T2)	4	12	66.7%
wrc	`dmb ishld` (T1, T2)	2	6	66.7%
mp	`dmb ishst` (T0); `ishld` (T1)	3	8	62.5%
s	`dmb ishst` (T0); `ishld` (T1)	3	8	62.5%
isa2	`dmb ishld`	2	4	50.0%
lb	`dmb ishld` (T0, T1)	2	4	50.0%
r	`dmb ishld`	2	4	50.0%

Table 14: Safety rates by pattern category.

Category	Patterns	Safe/Total	Rate
Fenced	10	99/100	99.0%
Asymmetric fence	6	42/60	70.0%
GPU scope	18	125/180	69.4%
Coherence	4	20/40	50.0%
Multi-thread	22	86/220	39.1%
Dependency	6	12/60	20.0%
Basic ordering	9	9/90	10.0%
Mutex	2	0/20	0.0%

I Performance Details

The sub-second analysis time enables integration into CI pipelines (checking all 570 pairs on every commit) and interactive developer workflows.

J Portability Checking Algorithm: Formal Specification

Algorithm 1 Cross-Architecture Portability Check

Require: Pattern $P = (ops, addrs, forbidden)$, target architecture A
Ensure: `PortabilityResult`: safe, fence recommendation

```
graph ← build_execution_graph(P)
stores_by_addr ← group_stores(P.ops, P.addrs)
rf_assignments ← enumerate_reads_from(graph)
co_assignments ← enumerate_coherence(stores_by_addr)
for each (rf, co) ∈ rf_assignments × co_assignments do
    E ← (graph, rf, co)
    if check_consistent(E, A) then
        outcome ← extract_outcome(E)
        if outcome = P.forbidden then
            return PortabilityResult(safe = False, fix = compute_fences(P, A))
        end if
    end if
end for
return PortabilityResult(safe = True)
```

Table 15: Per-pattern analysis time (10 architectures, mean of 100 runs).

Pattern	Threads	Mean (ms)	Stdev (ms)
mp	2	0.63	0.10
sb	2	0.64	0.10
lb	2	0.66	0.11
iriw	4	3.34	0.50
2+2w	3	6.01	0.87
rwc	3	1.38	0.20
wrc	3	1.57	0.24
mp_fence	2	2.14	0.35
sb_fence	2	1.17	0.19
isa2	3	1.46	0.22
Full 570 pairs (57 patterns)		135	31

Algorithm 2 Architecture Consistency Check

Require: Execution $E = (\mathit{graph}, \mathit{rf}, \mathit{co})$, architecture A
Ensure: Boolean: is E consistent under A?

```
ppo ← preserved_program_order(A, E.graph) {Filter po→ by A's relaxations}
for each fence f in E do
    ppo ← ppo ∪ fence_ordering(f, A) {Fences restore ordering}
end for
for each dependency d in E do
    if A.preserves_dependency(d.type) then
        ppo ← ppo ∪ {d.src → d.dst}
    end if
end for
ghb ← ppo ∪ rf_e ∪ co ∪ fr {Global happens-before}
return is_acyclic(ghb)
```

K Comparison with herd7: Qualitative Analysis

While Litmus∞ and `herd7` agree on all 50 compared CPU test-model pairs, the tools serve different purposes:

- **herd7**: Accepts `.litmus` files and `.cat` model specifications. Can simulate arbitrary user-defined models. Outputs allowed/forbidden for each outcome. Does not provide portability queries or fence recommendations.
- **Litmus∞**: Focuses on the *portability question*: "Is pattern P safe from A to B?" Includes GPU scoped models, per-thread fence recommendations, and scope mismatch detection. Does not support custom `.cat` models.

Complementary use: A practitioner can use Litmus∞ for rapid portability screening (570 pairs in 135 ms) and then use `herd7` for detailed investigation of specific failures with custom model definitions.

L Reproducibility

All experimental results in this paper are generated by the following commands:

Listing 3: Reproducing all experiments

```
# Full experiment suite (generates paper_results/)
python3 run_all_experiments.py

```

```bash
# Case studies with LLM integration
source ~/.bashrc  # set OPENAI_API_KEY
python3 run_case_studies.py

# Individual experiments
python3 portcheck.py --analyze-all
python3 portcheck.py --analyze-all --json
python3 fence_cost.py
```

All data files are saved in `paper_results/` and can be verified against the numbers reported in this paper.